

1 BEAVER: Practical Deterministic Verification of LLMs via 2 Frontier Exploration 3

4 ANONYMOUS AUTHOR(S)
5

6 As large language models (LLMs) transition from research prototypes to production systems, practitioners
7 often need reliable methods to verify that model outputs satisfy required constraints. While sampling-based
8 estimates provide an intuition of model behavior, they offer no formal guarantees. We present **BEAVER**,
9 the first practical framework for computing deterministic, sound probability bounds on LLM constraint
10 satisfaction. Given any prefix-closed semantic constraint, **BEAVER** systematically explores the generation
11 space using novel *token trie* and *frontier* data structures, maintaining provably sound bounds at every iteration.
12 We formalize the verification problem, prove soundness of our approach, and evaluate **BEAVER** on correctness
13 verification (GSM-Symbolic) and privacy verification (Enron email leakage) tasks across multiple state-of-the-
14 art LLMs. **BEAVER** achieves 2 to 8 times tighter probability bounds compared to baseline methods under
15 identical computational budgets, enabling precise characterization of model behavior that loose bounds cannot
16 provide.

17 1 Introduction

18 Large language models have demonstrated remarkable capabilities across diverse domains, from
19 engaging in complex conversations [1, 37] to driving scientific discovery [12, 26] and advancing
20 mathematical reasoning [13, 31]. As these models increasingly transition from research prototypes
21 to production systems, ensuring their reliability and safety has become paramount for real-world
22 deployment. Like classifiers in vision domains, there are a variety of risks associated with LLMs (e.g.,
23 privacy, safety) that must be evaluated before their real-world deployment. While there has been
24 a lot of work on deterministically verifying the safety properties of vision classifiers [39, 43, 47],
25 providing any type of deterministic guarantees on LLMs are generally considered infeasible due to
26 their enormous sizes. As a result, practitioners resort to either ad-hoc approaches based on bench-
27 marking [30], red-teaming [36], and adversarial attacks [58] or settle for statistical guarantees [11].

28 In this work, we demonstrate that deterministic verification of LLMs is both possible and practical.
29 Unlike traditional neural networks, LLMs are auto-regressive models that induce a distribution
30 over output sequences rather than producing a single deterministic output. At each generation step,
31 the model outputs a probability distribution over its vocabulary, conditioned on the prompt and
32 previously generated tokens. Overall, the LLM does not produce a single output, instead it induces
33 a probability distribution on the set of all possible output sequences for a given prompt. This
34 probabilistic nature fundamentally changes the verification problem. Rather than checking whether
35 a property holds on all outputs, we must compute the probability that the output distribution
36 satisfies a given constraint. This paper tackles the first foundational step: we provide a method to
37 compute deterministic, sound bounds on constraint satisfaction probability for a single prompt.

38 We consider verifying LLMs with respect to prefix-closed semantic constraints on their outputs,
39 which are a rich class of decidable predicates where if a prefix violates a constraint any continuation
40 is also violating. These predicates can capture properties such as correctness, privacy, and safety
41 (as shown in our experiments). In order to compute the models constraint-satisfaction probability,
42 we must find the total probability mass of all model responses that satisfy our constraints. However,
43 computing this probability exactly is intractable. With vocabulary sizes exceeding one hundred
44 thousand tokens and even moderate sequence lengths, the output space grows exponentially, a
45 combinatorial explosion that precludes exhaustive enumeration.

46 Because of the differences in how LLMs work, we cannot directly build on top of traditional
47 techniques based on abstract interpretation [42] or SMT solvers [27]. These approaches aim to
48 certify properties of a single pass from inputs to outputs. In contrast, LLMs compute using an
49

50 auto-regressive process based on multiple forward passes that combines high-dimensional continuous
 51 computations with discrete, algorithm-dependent decoding steps. Modeling this mixture of
 52 probabilistic choice, sequential unrolling, and decoding logic falls outside the expressiveness and
 53 scalability of current symbolic verification frameworks, which would either not scale to LLM-sized
 54 architectures or yield vacuous over-approximations. Therefore, new verification principles are
 55 needed to reason soundly about the probabilistic semantics of LLMs.

56 We present **BEAVER**, a novel framework that computes provably sound probability bounds
 57 for LLM constraint-satisfaction through systematic exploration of the generation space. Our key
 58 insight is that for prefix-closed semantic constraints, we can aggressively prune the search space by
 59 detecting and discarding constraint violations as soon as they occur. **BEAVER** maintains two novel
 60 data structures: ① A token trie that explicitly tracks all explored constraint-satisfying prefixes
 61 along with their probabilities, and ② A frontier representing complete and incomplete sequences
 62 used for bound computation. At each step, **BEAVER** selects an incomplete token sequence from
 63 the frontier, performs a single model forward pass to obtain its next-token distribution, adds all
 64 constraint-satisfying continuations to the token trie, and updates sound lower and upper bounds on
 65 the target probability. By maintaining these monotonically tightening bounds throughout execution,
 66 **BEAVER** provides anytime guarantees, that it at any point, practitioners can terminate with sound
 67 probability intervals.

68
 69 **Main Contributions.** Our work provides the first practical framework for soundly computing
 70 deterministic probability bounds on LLM constraint satisfaction:

- 71 • **Formal Framework:** We formalize the LLM deterministic verification problem as computing
 72 probability bounds over constraint-satisfying generations and present novel token
 73 trie and frontier data structures defined over the LLM generation that enable sound bound
 74 computation.
- 75 • **BEAVER Algorithm:** We present our branch-and-bound verification algorithm with formal
 76 soundness proofs, demonstrating that our bounds are valid at every iteration and converge
 77 toward the true probability with additional computation.
- 78 • **Empirical Validation:** We evaluate **BEAVER** on two critical verification tasks: correctness
 79 verification using the GSM-Symbolic mathematical reasoning benchmark [31] and privacy
 80 verification using an email leakage task [33], across multiple state-of-the-art LLMs. Our
 81 results show that **BEAVER** achieves 2-8 times tighter probability bounds compared to
 82 rejection sampling baselines under identical computational budgets.

84 85 2 Background

86 In this section, we provide the relevant background on language models, formal language grammar
 87 and semantic constraints.

88 Notation

91 We use Σ to denote an alphabet (a finite set of symbols), and Σ^* to denote the set of all finite strings
 92 over Σ , including the empty string ϵ . Any vector or sequence of elements is written using bold
 93 characters \mathbf{a} . Additionally, we use $\mathbf{a} \preceq \mathbf{b}$ to denote that \mathbf{a} is a **prefix** of \mathbf{b} , and $\mathbf{a} \prec \mathbf{b}$ to denote a
 94 **strict prefix** relation. The \cdot operator is used to denote concatenation for sequences and elements
 95 to sequences, that is $\mathbf{a} = \mathbf{b} \cdot \mathbf{c}$ implies that \mathbf{a} is the concatenation of 2 sequences \mathbf{b} and \mathbf{c} , and
 96 $\mathbf{d} = \mathbf{e} \cdot f$ is the concatenation of \mathbf{e} with element f . For any natural number $i \in \mathbb{N}$, $[i]$ denotes the
 97 set $\{j \mid 1 \leq j \leq i\}$.

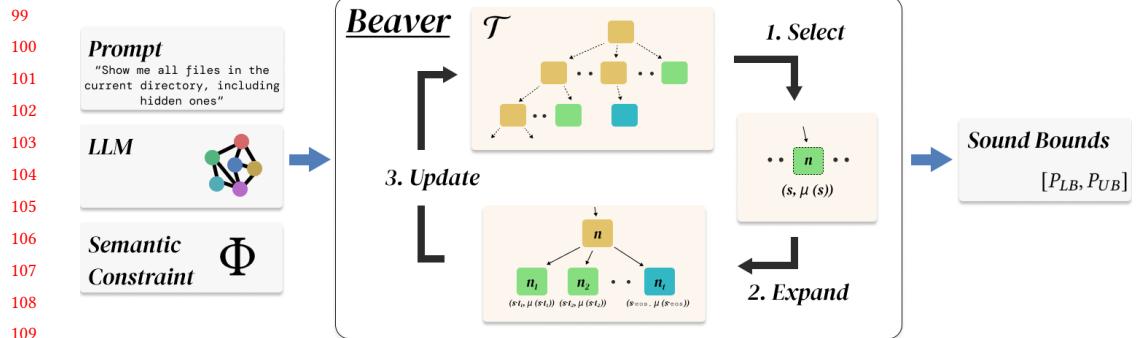


Fig. 1. **BEAVER** workflow for computing sound probability bounds. Given a prompt, language model, and a prefix-closed semantic constraint, **BEAVER** iteratively: (1) selects an incomplete leaf from the frontier, (2) expands it by querying the model and adding valid continuations to the token trie, and (3) updates the sound probability bounds $[P_{LB}, P_{UB}]$ based on the new frontier state.

2.1 Language Models

Language models M operate on vocabulary of tokens $\Sigma \subseteq V \subseteq \Sigma^*$. A tokenizer $\tau : \Sigma^* \rightarrow (V \setminus \{\langle \text{eos} \rangle\})^*$ takes input any string $\mathbf{p}_u \in \Sigma^*$, commonly called the user prompt, and convert it into a sequence of tokens $\mathbf{p} = \tau(\mathbf{p}_u) = t_1 \cdot t_2 \cdots t_k$ where $t_i \in (V \setminus \{\langle \text{eos} \rangle\})$. This sequence of tokens is taken as input by M , which returns a vector of real numbers of the size $|V|$, referred to as logits \mathbf{z} . We apply the softmax function $\text{softmax}(z_i) = e^{z_i} / \sum_j e^{z_j}$ on logits to get a *probability distribution* $P_M(\cdot | \mathbf{p})$ over V , used for predicting the next token in the sequence of tokens given input. This process of generating $P_M(\cdot | \mathbf{p})$ for the next token following prompt \mathbf{p} is referred to as a *forward pass*.

Decoding. After a forward pass on prompt \mathbf{p} , a token $t \in V$ is selected based on the probability distribution $P_M(\cdot | \mathbf{p})$. This token is appended to the end of the input prompt, and fed back to the language model to get the following token distribution $P_M(\cdot | \mathbf{p} \cdot t)$. This step is repeated multiple times to get a sequence of tokens following the prompt \mathbf{p} . This iterative process stops when a certain $\langle \text{eos} \rangle \in V$ (end-of-sequence) token is sampled. The resulting sequence of tokens \mathbf{r} following the prompt \mathbf{p} is called a *response*. Each response \mathbf{r} is a sequence of tokens of the form $\mathbf{r} = \{t_1 \cdot t_2 \cdots t_n \cdot \langle \text{eos} \rangle\} \mid t_i \in V \setminus \{\langle \text{eos} \rangle\}$. The generated list of tokens can then be de-tokenized by the tokenizer to give a final output response string.

Let $P_M(\cdot | x_1 \cdot x_2 \cdots x_{p-1})$ denote the probability distribution over the vocabulary V produced by a language model M , conditioned on the token sequence $x_1 \cdot x_2 \cdots x_{p-1}$. We define $\mu(\mathbf{s}_n)$ as the model's probability of generating token sequence $\mathbf{s}_n = \{t_1 \cdot t_2 \cdots t_n\}$ given input prompt \mathbf{p} below.

$$\mu(\mathbf{s}_n) = \prod_{i=1}^n P_M(t_i | \mathbf{p}_{i-1}) \quad \text{where } \mathbf{p}_0 = \mathbf{p} \text{ and } \mathbf{p}_i = \mathbf{p}_{i-1} \cdot t_i \text{ for all } i \in [n] \quad (1)$$

Various token selection strategies to, referred to as *decoding strategies*, have been explored in the literature for different objectives such as maximum likelihood or diversity. We cover some of them in Appendix A. Current language models are capable of learning sufficient probability distributions to be able to answer questions and solve tasks with extensive training on natural and programming languages. However, they fail to learn complex tasks, or due to the probabilistic nature of their response generation, are not able to consistently follow formal language rules.

148 *Rejection Sampling.* In order to get responses from language model that satisfy a given formal
 149 / semantic constraint, various strategies exist. One of the simplest and most common strategy is
 150 to iteratively sample responses from the model till one correctly satisfies the given constraint.
 151 This method of repeated sampling for constraint satisfaction is called *rejection sampling*. While
 152 rejection sampling generates probable responses, for restrictive or complex constraints, this ap-
 153 proach may require generating a very large number of samples before finding a valid one, making
 154 it computationally inefficient.

155 2.2 Semantic constraints

156 Beyond basic token generation, we often need to verify that language model outputs satisfy specific
 157 requirements. These requirements may include syntactic validity (e.g., well-formed JSON), security
 158 properties (e.g., no dangerous operations), functional correctness (e.g., passes test cases), or any
 159 combination thereof. We formalize such requirements as *semantic constraints*.

160 *Definition 2.1 (Semantic Constraint).* A semantic constraint is a decidable predicate $\Phi : V^* \rightarrow \{\top, \perp\}$ over token sequences. For a sequence $s \in V^*$, we say that s satisfies constraint Φ , written
 161 $s \models \Phi$, if and only if $\Phi(s) = \top$

162 We require that Φ be decidable. There must exist an algorithm that, given any finite token
 163 sequence $s \in V^*$, determines whether $\Phi(s)$ in finite time. This requirement is essential for practical
 164 verification.

165 The above definition of semantic constraints allows a wide variety of specifications to be encoded
 166 as a semantic constraint. For example, for the regex R : “`^\d{4}-\d{2}-\d{2}$`”, which is the regex
 167 constraint for valid date in YYYY-MM-DD format, one can define semantic constraint Φ_R as one which
 168 checks a token sequence $s \in V^*$ satisfies the given regex constraint. that is $s \models \Phi_R \Leftrightarrow s \in \mathcal{L}(R)$.
 169 Another example constraint could be Φ_{toxic} which checks if a token sequence has some keywords
 170 that indicate toxicity.

171 A critical property for semantic constraints, is *prefix-closure*.

172 *Definition 2.2 (Prefix-closed semantic constraints).* A semantic constraint $\Phi : V^* \rightarrow \{\top, \perp\}$ is
 173 *prefix-closed* if for all token sequences, if s satisfies the constraint Φ , then any prefix $s' \preceq s$ also
 174 satisfies the constraint Φ . That is,

$$175 \quad \forall s, s' \in V^*, s \models \Phi \wedge s' \preceq s \implies s' \models \Phi$$

176 Equivalently, if any prefix s' violates Φ , then all extensions of s' also violate Φ :

$$177 \quad \forall s, s' \in V^*, s' \not\models \Phi \wedge s' \preceq s \implies s \not\models \Phi$$

178 This property is importantly crucial for our proposed approach, since it enables us to check if a
 179 given subset of token sequences with the same prefix violate the given semantic constraint.

180 Many natural constraints are *not* prefix-closed. Consider the date regex constraint Φ_R defined
 181 above. This constraint is *not* prefix-closed as for $s = "2024-10-15"$ and $s' = "2024"$, $s' \not\models \Phi_R$ but
 182 $s \models \Phi_R$. However, we can make a new constraint Φ_{R_p} which is a prefix-closed variant of Φ_R .

183 *Definition 2.3 (Complete Token Sequences).* The complete token sequences C denoted by the set
 184 of strings $C = (V \setminus \langle \text{eos} \rangle)^* \langle \text{eos} \rangle$. This essentially captures all valid token sequences the LLM M can
 185 produce as M always stops autoressive generation post the $\langle \text{eos} \rangle$ token generation.

186 Next, we define the verification problem.

197 *Definition 2.4 (verification problem).* Given an input LLM M , tokenized input prompt \mathbf{p} , and a
198 semantic constraint Φ , we want to find the total probability P of strings $\mathbf{s} \in C$ satisfying Φ , where
199 these strings \mathbf{s} are drawn from the LLM predicted distribution $P_M(\cdot | \mathbf{p})$ on tokenized input \mathbf{p} .

200 Formally the value of P is given by the following equation
201

$$\small P = \sum_{\mathbf{s}_i \in C} \mu(\mathbf{s}_i) * \mathbb{1}[\mathbf{s}_i \models \Phi] \quad (2)$$

202 Computing P exactly requires enumerating all responses in C , checking which satisfy Φ , and
203 summing their probabilities. Just to give an idea of size of the set C even if we restrict ourselves
204 to only token sequences of length $L = 6$, with a vocabulary $|V| = 15$, $O(C) = |V|^{L-1}$ which is equal
205 to $15^5 = 759375$ sequences. This becomes further intractable for realistic vocabularies ($|V| \sim 50000$).
206 Instead, to achieve practical runtime, we obtain sound interval bound $P_{LB} \leq P \leq P_{UB}$ and we
207 iteratively tighten the bounds while maintaining soundness over all iteration. In the next sections,
208 we develop an approach that computes these bounds incrementally without enumerating over C .
209

210 3 Overview

211 Figure 1 illustrates the core idea behind **BEAVER** framework. Given a language model M and
212 a prefix-closed semantic constraint Φ , our method computes provably sound probability bounds
213 P_{UB}, P_{LB} of the model generating constraint-satisfying response for a given input.

214 Our key insight is that we can track partial sequences and prune constraint violations early.
215 Our algorithm maintains a novel *Token Trie* which tracks all partial constraint-satisfying token
216 sequences along with their probabilities and a *frontier* to track valid incomplete sequences. Unlike
217 a baseline approach like rejection sampling, which wastes forward passes on duplicate samples
218 and examines entire sequences even when early tokens already violate the constraint, our frontier-
219 based approach (1) tracks an explicit search state (*frontier*) to avoid redundant work, (2) leverages
220 prefix-closure property of the constraints to prune entire subsets of possible generations and (3)
221 progressively refines bounds by exploration of high-probability prefixes. This enables our method
222 to achieve much tighter bounds than baseline approaches, making formal verification of LLM
223 behavior practical even under computational budgets.

224 We illustrate our approach through a concrete toy example to illustrate our frontier-based
225 verification algorithm. This section builds intuition for the formal treatment in Section 4. We begin
226 with a running example in Section 3.1 that demonstrates the need for computing bounds on the
227 specified constraint.

228 We then examine the baseline method and its inefficiencies in Section 3.2, before introducing
229 and providing a walkthrough of our **BEAVER** algorithm in Section 3.3 and 3.4.

230 3.1 Illustrative Example

231 *3.1.1 The Task.* We consider a language model M tasked with generating bash commands in
232 response to natural language queries. It has a simplified vocabulary V of 16 tokens and can generate
233 a response of maximum length $L = 5$.

$$\small V = \{\text{ls, rm, cat, chmod, cd, echo, -la, -rf, -R, -l, ., /home, /tmp, /etc/passwd, ~, <eos>}\}$$

234 For a given prompt \mathbf{p} , each response $\mathbf{r} \in C$ is a sequence of atmost L tokens from V and ends
235 with the $\langle \text{eos} \rangle$ token, i.e. $\mathbf{r} = \{t_1 \cdot t_2 \cdots t_n \cdot \langle \text{eos} \rangle \mid n < L, t_i \in V\}$ (Definition 2.3). Following section
236 2.1, the probability of generating response r is defined in Eq. 1.

237 For the prompt \mathbf{p} : “Show me all files in the current directory including hidden ones”, The expected
238 safe output is “`ls -al`”. However the model’s vocabulary also permits it to generate unsafe
239 commands such as “`rm -rf /home`”.

Sequence \mathbf{s}	Probability $\mu(s)$	Sample count
[ls-al. $\langle eos \rangle$]	0.21	4
[ls-al $\langle eos \rangle$]	0.168	2
[ls. $\langle eos \rangle$]	0.07	3
[rm-rf $\langle eos \rangle$]	0.07	1

Table 1. Sequences sampled with rejection sampling for the safe bash command example.

Our goal is to find the probability of the model to generate a safe command.

3.1.2 Safety Constraint Φ . In order to formally define safe / unsafe commands, we define a safety specification Φ which requires:

- No deletion operations (rm commands).
- No accesses to sensitive system files (/etc/passwd)
- No permission modifications (chmod commands)

We define $\Phi : V^* \rightarrow \{\top, \perp\}$ as a semantic constraint (refer to section 2.2) which is a predicate over token sequences. Token sequence $\mathbf{s} \models \Phi$ if and only if \mathbf{s} satisfies all the safety requirements listed above. Formally, we check for Φ in token sequence \mathbf{s} as :

$$\Phi \models \mathbf{s} \Leftrightarrow \neg(\text{rm} \in \mathbf{s}) \wedge \neg(\text{/etc/passwd} \in \mathbf{s}) \wedge \neg(\text{chmod} \in \mathbf{s})$$

Crucially, safety constraint Φ is *prefix-closed*. While generating a response, if the first token generated by the model is “rm”, any continuation from this token will also violate our safety constraint.

3.2 Provable bounds using rejection sampling

3.2.1 Baseline rejection sampling. We wish to compute sound lower and upper bounds $[P_{LB}, P_{UB}]$ on the probability of generating a constraint-satisfying response from the model. A naive approach to compute these bounds is through *rejection sampling*. In rejection sampling, we iteratively sample complete sequences along with their probabilities $(\mathbf{s}, \mu(\mathbf{s}))$ from the model. We start by maximally setting our lower bound $P_{LB} = 0.0$ and our upper bound $P_{UB} = 1.0$. For each sampled sequence s , if $\mathbf{s} \models \Phi$, we increase our lower bound by adding $\mu(\mathbf{s})$ to P_{LB} . Else if $\mathbf{s} \not\models \Phi$, we tighten the upper bound by subtracting $\mu(\mathbf{s})$ from P_{UB} . The detailed algorithm for bound calculation using rejection sampling can be found in Appendix B. This approach provides sound bounds since only probabilities of safe responses contribute to P_{LB} , while unsafe responses are removed from P_{UB} , maintaining $P_{LB} \leq P \leq P_{UB}$ at all iterations.

3.2.2 Walkthrough on the bash example. Consider our above example with vocabulary $|V| = 15$ (excluding $\langle eos \rangle$) and maximum length $L = 5$. We initialize $P_{UB} = 1.0$ and $P_{LB} = 0.0$. Suppose we sample 10 sequences from the model with rejection sampling. Table 1 presents the sequences sampled along with their probabilities and frequencies. As shown in the table, only 4 novel sequences were obtained despite 34 forward passes. Since sequences [ls-al. $\langle eos \rangle$], [ls-al $\langle eos \rangle$], and [ls. $\langle eos \rangle$] satisfy the safety constraint Φ , their sequence probabilities are added to the lower bound. $P_{LB} = 0.21 + 0.168 + 0.07 = 0.378$. Conversely, since [rm-rf $\langle eos \rangle$] violates Φ and has probability 0.07, $P_{UB} = 1 - 0.07 = 0.93$. The resultant bounds $[P_{LB}, P_{UB}] = [0.378, 0.93]$ have gap 0.552.

295 3.2.3 *Inefficiencies of Rejection Sampling.* The walkthrough above highlights several fundamental
296 inefficiencies that make naive rejection sampling impractical for computing tight constraint-
297 satisfaction bounds.

298 First, duplicate sampling quickly dominates the computation. As more sequences are drawn, the
299 probability of resampling high-probability sequences grows rapidly. In our bash example (Table 1),
300 only three distinct sequences are discovered, while seven of the ten samples were duplicates that
301 provided no new information. To avoid duplicates, we need to keep track of all expanded sequences
302 and only sample those which we have not yet explored.

303 Second, rejection sampling does not fully exploit the prefix-closure property of our safety
304 constraint Φ . It continues generation until $\langle \text{eos} \rangle$ token is generated and updates the bounds only
305 at the end of the sequence, wasting up to $O(L)$ extra model forward passes per unsafe sample. To
306 avoid wasting forward passes, we need to prune prefix \mathbf{x} as soon it violates the constraint and to
307 subtract its probability $\mu(\mathbf{x})$ from the upper bound.

308 Taken together, these observations suggest that addressing duplicate sampling and exploiting
309 prefix-closure requires efficiently tracking partial sequences. In the next section, we introduce a
310 tree data structure over partial sequences that forms the basis of our **BEAVER** algorithm.

312 3.3 BEAVER Data Structures

313 We now present the core intuition behind our **BEAVER** approach. **BEAVER** explicitly maintains
314 the set of all partial sequences that satisfy the semantic constraint. It then uses this set to compute
315 probability bounds for constraint satisfaction. **BEAVER** exploits the prefix-closure property of
316 Φ (Definition 2.2) and rejects any partial generation that violates Φ . **BEAVER** tracks all these
317 sequences using a novel trie data structure called the *token trie* \mathcal{T} . Figure 2 illustrates the token trie
318 structure and shows how **BEAVER** updates it in our bash command example.

320 **Trie structure.** Each node in \mathcal{T} corresponds to a sequence \mathbf{s} formed by concatenating all tokens
321 along the path from the root to that node, and its sequence probability $\mu(\mathbf{s})$ is the product of edge
322 probabilities along this path. The root of the trie represents the empty sequence ϵ . Each edge in
323 the token trie is labeled with (1) a token $t \in V$ and (2) the conditional probability $P_M(t | \mathbf{p} \cdot \mathbf{s})$ of
324 generating t given prompt \mathbf{p} and sequence \mathbf{s} corresponding to the parent node. We use the notation
325 $n[\mathbf{x}]$ for node n in the token trie, which represents sequence \mathbf{x} .

326 In Figure 2, the root node n_0 has three children: n_1 reached by token `ls` with probability 0.6, n_2
327 reached by token `echo` with probability 0.2, and n_3 reached by token $\langle \text{eos} \rangle$ with probability 0.01.
328 Tokens such as `rm` and `chmod` are *not* added as children of n_0 because they immediately violate Φ .

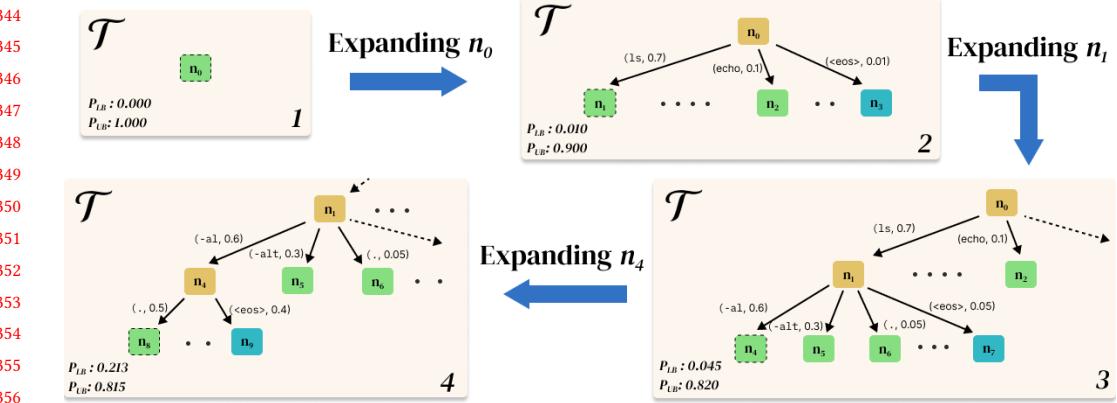
329 A leaf node is *complete* if and only if its incoming edge token is $\langle \text{eos} \rangle$, indicating that the model
330 has finished generating the sequence (turquoise nodes in the figure). For instance, the leaf node n_6 ,
331 representing the sequence `echo · <eos>`, is complete because it ends in $\langle \text{eos} \rangle$. Leaf nodes that are
332 not complete are *incomplete* and are eligible for expansion (colored in green in the figure).

333 **Frontier.** The collection of all leaf nodes which in a token trie \mathcal{T} is called a *frontier* Ψ . The
334 frontier is the set of leaf nodes for a given iteration of **BEAVER**. Ψ splits into two sets: Ψ_c , which
335 is the set of complete leaves, and Ψ_i , which is the set of incomplete leaves.

337 3.4 BEAVER Walkthrough

339 Initially, the trie \mathcal{T} starts with just the root node corresponding to the empty sequence ϵ . $[P_{LB}, P_{UB}] =$
340 $[0, 1]$. **BEAVER** grows \mathcal{T} through iterative expansion of incomplete leaves from Ψ_i . Each iteration
341 consists of three steps: **Select**, **Expand**, and **Update**.

342 **Selection:** **BEAVER** first selects an incomplete leaf $u \in \Psi_i$ from the frontier.



357
358 Fig. 2. Evolution of the token trie \mathcal{T} through four iterations of **BEAVER** on the bash command safety
359 constraint. Starting from the empty trie, **BEAVER** expands nodes n_0 , n_1 , and n_4 in sequence. Green nodes
360 indicate incomplete sequences eligible for expansion, turquoise nodes indicate complete sequences (ending
361 in $\langle eos \rangle$). Probability bounds tighten from $[0.01, 0.9]$ after iteration 1, to $[0.213, 0.815]$ after iteration 3, to
362 $[0.7, 0.8]$ after iteration 10. Low probability sequence nodes omitted for brevity.

363
364
365 **Expansion:** **BEAVER** queries the model for the probability distribution $P_M(\cdot | p \cdot x)$ over
366 vocabulary V . For each token $t \in V$ such that $x \cdot t \models \Phi$, **BEAVER** adds a new child node to the
367 node u corresponding to $x \cdot t$ with an edge with the label $(t, P_M(t | p \cdot x))$. $n[x \cdot t]$ is complete
368 if $t = \langle eos \rangle$ and incomplete otherwise. This turns the former incomplete leaf u into an internal
369 node. Updating the trie \mathcal{T} to \mathcal{T}' correspondingly updates the frontier Ψ to Ψ' . Specifically, $n[x]$
370 is removed from Ψ_i as it is expanded to new valid sequences. Child node $n[x \cdot \langle eos \rangle]$ is added to
371 Ψ_c , while all other new nodes corresponding to constraint-satisfying continuations $n[x \cdot t]$ where
372 $t \in V \setminus \{\langle eos \rangle\}$ are added to Ψ_i .

373 **Updating Bounds:** **BEAVER** uses the updated frontier Ψ' to compute probability bounds P_{LB}
374 and P_{UB} on P . The lower bound $P_{LB}[\Psi']$ sums the probabilities of all complete sequences in Ψ'_c ,
375 representing sequences we have certified that satisfy Φ . The upper bound $P_{UB}[\Psi']$ is computed
376 based on both complete and incomplete sequences, treating each incomplete sequence $x \in \Psi'_i$ as if
377 all of its continuations satisfy Φ and thus contributing its full probability mass $\mu(x)$. In Section 4.4,
378 we show that these bounds are sound and monotonic.

379
380 3.4.1 **Walkthrough for bash command example.** Consider the example shown in Figure 2,
381 which shows three iterations of **BEAVER** for the safe bash command task. In this example, at each
382 iteration, **BEAVER** expands the incomplete node with the highest sequence probability. We describe
383 this selection strategy in detail in Section 4.2. Note: Figure 2 only shows nodes corresponding to
384 sequences with non-trivial probabilities. Other nodes with lower probabilities are omitted from the
385 figure for brevity, but are still included in the bound computation.

386 Starting from the empty trie \mathcal{T} , after the iteration 1 where root n_0 is expanded, our frontier is
387 updated to now include nodes corresponding to incomplete sequences $\{(1s, 0.7), (\text{echo}, 0.1), \dots\}$ and nodes corresponding to complete sequence $\{(\langle eos \rangle, 0.01)\}$. Crucially, **BEAVER** prunes out prefixes $\{\text{rm}, \text{chmod}, \dots\}$, whose sequence probabilities sum up to 0.1,
388 which violate our safety constraint Φ . Thus, the probability bounds after iteration 1 are $P_{LB} = 0.01$,
389 $P_{UB} = 0.9$.
390
391

393 After iteration 2 where node n_1 is expanded, the frontier is

$$394 \quad \Psi_c = \{n[\langle \text{eos} \rangle], n[1s \cdot \langle \text{eos} \rangle]\}$$

$$395 \quad \Psi_i = \{n[1s \cdot -al], n[1s \cdot -alt], n[echo], n[1s \cdot .], \dots\}.$$

396 Thus, the probability bounds after iteration 2 are:

$$397 \quad P_{LB}[\Psi] = 0.045, \quad P_{UB}[\Psi] = 0.82$$

400 After ten iterations, **BEAVER** finds high probability valid completed sequences $[1s \cdot -al \cdot . \langle \text{eos} \rangle]$,
 401 $[1s \cdot -al \langle \text{eos} \rangle]$, and $[1s \cdot -alt \langle \text{eos} \rangle]$, increasing the lower bound, while decreasing the upper bound
 402 by pruning out more invalid prefixes. The probability bounds after 10 iterations are:
 403

$$404 \quad P_{LB}[\Psi'] = 0.7, \quad P_{UB}[\Psi'] = 0.8$$

405 The running example highlights several crucial aspects of **BEAVER**. Firstly, by maintaining
 406 a trie and frontier over possible sequences that satisfy the constraint Φ , **BEAVER** exploits the
 407 prefix-closure property of Φ and prunes out thousands of violating sequences early on. In our
 408 example, tokens such as `rm` and `chmod` are discarded at iteration 1, which rules out a large mass of
 409 unsafe sequences. On the other hand, rejection sampling only discovers a violation after generating
 410 a full sequence such as $[rm \cdot -rf \cdot /home \cdot \langle \text{eos} \rangle]$. Furthermore, **BEAVER**'s bounds tighten rapidly.
 411 The gap between upper and lower bound achieved after 10 model forward passes is 0.1, while
 412 rejection sampling only manages to reduce the gap to 0.552 despite over three times as many model
 413 forward passes.
 414

415 4 LLM Verification with Branch and Bound

416 In this section, we introduce the relevant data structures (Section 4.1) and outline the **BEAVER**
 417 algorithm (Section 4.2), which incrementally updates the lower bound P_{LB} and the upper bound
 418 P_{UB} while maintaining the soundness condition $P_{LB} \leq P \leq P_{UB}$ at each step. The pseudocode is
 419 provided in Section 4, and we formally prove the soundness of **BEAVER** in Section 4.4.

420 4.1 Incremental bound computation via Frontiers

421 For efficiently computing the probability bounds, we only need to track the set of possible sequences
 422 that satisfy the constraint and use this set to compute the bounds. This approach allows us to
 423 exploit the prefix-closure property of Φ (Definition 2.2) and to reject early any sequences that
 424 already violate Φ . To this end, we modify the trie data structure [15] (referred to as the token trie
 425 \mathcal{T}) to track all possible constraint-satisfying sequence generations produced by the model for a
 426 given prompt \mathbf{p} . We then define a frontier on this trie, representing the current set of valid partial
 427 sequences (those not ending with the $\langle \text{eos} \rangle$ token) and completed sequences (Definition 2.3). Next,
 428 we provide necessary definition and update rules for \mathcal{T} .
 429

430 *Definition 4.1 (Token Trie).* We model LLM sequence generation as incrementally constructing a
 431 trie (prefix-tree) \mathcal{T} over token sequences that satisfy constraint Φ . By the prefix-closure property
 432 of Ψ (Definition 2.2), any continuation of a constraint-violating sequence also violates Ψ . Therefore,
 433 we only track constraint-satisfying sequences in \mathcal{T} .
 434

435 **Trie Structure:** The root node represents the empty sequence ϵ . We representation of edges
 436 and nodes of \mathcal{T} below

- 437 • **Edge:** Each edge is labeled with: 1) a token $t \in V$ and 2) the conditional probability $P_M(t \mid \mathbf{p} \cdot \mathbf{s})$
 438 of generating that token given the prompt \mathbf{p} and the sequence \mathbf{s} of the parent node.

- 442 • **Node:** Each node's label contains the token sequence \mathbf{s} obtained by concatenating edge token
 443 labels along the path from the root to that node and the sequence probability $\mu(\mathbf{s})$. We use $n[\mathbf{s}]$
 444 to denote the node with token sequence \mathbf{s} .

445 The sequence probability $\mu(\mathbf{s})$ can be computed by multiplying the conditional probabilities
 446 along this path. All token sequences represented in \mathcal{T} satisfy the constraint Φ . Recall, the LLM
 447 stops generation after generating the $\langle \text{eos} \rangle$ token. Hence, we say a node is *complete* if its incoming
 448 edge is labeled $\langle \text{eos} \rangle$; otherwise, it is *incomplete*.
 449

450 **\mathcal{T} Update Strategy:** The trie is updated incrementally after each token generation. Let u be
 451 an incomplete leaf in the trie and \mathbf{x} be the corresponding label sequence. In an update $\mathcal{T} \xrightarrow{\mathbf{x}} \mathcal{T}'$, for
 452 each token $t \in V$, we add an edge from u to a new child node labeled with token t if and only if
 453 $\mathbf{x} \cdot t \models \Phi$. After the update, u is no longer a leaf node.
 454

455 *Definition 4.2 (Frontier).* We define the *frontier* Ψ as the set of all leaf nodes in trie \mathcal{T} . Ψ is split
 456 into two disjoint sets: Ψ_c (complete leaves) and Ψ_i (incomplete leaves) ($\Psi = \Psi_c \cup \Psi_i$). For a trie
 457 update $\mathcal{T} \xrightarrow{\mathbf{x}} \mathcal{T}'$, the corresponding update to the frontier $\Psi \xrightarrow{\mathbf{x}} \Psi'$ is defined as

$$458 \quad \Psi'_c = \Psi_c \cup n[\mathbf{x} \cdot \langle \text{eos} \rangle], \quad \Psi'_i = (\Psi_i \setminus n[\mathbf{x}]) \cup \{n[\mathbf{x} \cdot t] \mid t \in V, \mathbf{x} \cdot t \models \Phi\} \quad (3)$$

460 In other words, Ψ'_c is updated with the sequence completing \mathbf{x} with $\langle \text{eos} \rangle$. Ψ'_i is updated with all
 461 constraint-satisfying next-token continuations of \mathbf{x} .

462 *Note on terminology:* Throughout this section, when context is clear, we may refer to "*expanding*
 463 *frontier* Ψ " as a shorthand for "*expanding the trie whose frontier is Ψ* ".

464 **Incremental Update of \mathcal{T} :** Initially, \mathcal{T} has just the root node labelled by the empty sequence
 465 $n[\epsilon]$. Hence, $\Psi_i = \{\epsilon\}$ and $\Psi_c = \emptyset$. At each update step, we select some incomplete leaf node $n[\mathbf{x}]$
 466 with corresponding token sequence \mathbf{x} . We perform one forward pass of M to obtain $P_M(\cdot \mid \mathbf{p} \cdot \mathbf{x})$.
 467 For each token $t \in V$, we add an edge from $n[\mathbf{x}]$ to a new child node labeled with token t and
 468 its conditional probability $P_M(t \mid \mathbf{p} \cdot \mathbf{x})$ if and only if $\mathbf{x} \cdot t$ satisfies the constraint $(\mathbf{x} \cdot t \models \Phi)$.
 469 Hence, for the updated trie \mathcal{T}' , Ψ'_c is updated with $\Psi'_c = \Psi_c \cup n[\mathbf{x} \cdot \langle \text{eos} \rangle]$. Ψ'_i is updated with all
 470 constraint-satisfying next-token continuations of \mathbf{x} (see Eq. 3).

471 **Iterative sound bound computation:** We define P_{UB} and P_{LB} at any step based on the frontier
 472 state. $P_{UB}[\Psi]$ is written as the sum of probability of all sequences in frontier $\Psi = \Psi_i \cup \Psi_c$, while
 473 $P_{LB}[\Psi]$ is written as the sum of probability of all sequences in Ψ_c .

$$475 \quad P_{UB}[\Psi] = \sum_{\mathbf{s}_i \in \Psi_i \cup \Psi_c} \mu(\mathbf{s}_i), \quad P_{LB} = \sum_{\mathbf{s}_i \in \Psi_c} \mu(\mathbf{s}_i), \quad P_{UB}[\Psi] - P_{LB}[\Psi] = \sum_{\mathbf{s}_i \in \Psi_i} \mu(\mathbf{s}_i) \quad (4)$$

477 Intuitively, $P_{LB}[\Psi]$ represents the total probability mass of all completed sequences (sequences
 478 that end with $\langle \text{eos} \rangle$) that satisfy Φ , which are captured by sequences corresponding to the leaf
 479 nodes in Ψ_c . Meanwhile, $P_{UB}[\Psi]$ represents the total probability mass of all sequences (both
 480 incomplete and complete) that satisfy constraint Φ , captured by sequences corresponding to leaf
 481 nodes in $\Psi = \Psi_i \cup \Psi_c$. The difference between them, $P_{UB}[\Psi] - P_{LB}[\Psi]$ represents the uncertain
 482 probability mass, the set of incomplete sequences (corresponding to leaf nodes in Ψ_i) that might or
 483 might not lead to valid completions.

4.2 Greedy Heuristic for Frontier Expansion

486 To efficiently tighten the certified bounds
 487 (P_{LB}, P_{UB}) under a strict budget of δ forward
 488 passes, we view frontier expansion as a search
 489 process in which each transition improves the

490 **Algorithm 1:** SearchSequence – Max- μ
 strategy

491 **Input :**Frontier Ψ with sequences \mathbf{s}
 492 keyed by $\mu(\mathbf{s})$

493 **Output:** $\mathbf{s}^*, \mu(\mathbf{s}^*)$

494 **1 return** $\arg \max_{\mathbf{s}_i} \mu(\mathbf{s}_i)$

491 bounds by expanding one sequence in the frontier Ψ . Since only one forward pass is permitted
 492 per expansion, the effectiveness of the verifier
 493 depends critically on choosing the sequence
 494 that yields the largest reduction in the proba-
 495 bility gap. Although computing the optimal choice is intractable due to prohibitive cost constraints,
 496 we employ a practical, lightweight best-first heuristic, $Max\text{-}\mu$, which always expands the sequence
 497 with the highest path probability. Formally, the selected sequence is
 498

$$x^* = \arg \max_{x \in \Psi_i} \mu(x).$$

501 We also implement a probabilistic strategy $Sample\text{-}\mu$ that samples incomplete sequences from
 502 the Ψ_i proportionally to their path probabilities. Formally, the selection probability for incomplete
 503 sequence \mathbf{x} in $Sample\text{-}\mu$ is
 504

$$P(\mathbf{x}) = \mu(\mathbf{x}) / \sum_{\mathbf{x}' \in \Psi_i} \mu(\mathbf{x}')$$

505 This strategy trades determinism for stochastic exploration, potentially discovering diverse high-
 506 probability paths earlier in verification but sacrificing the guarantee of always expanding the most
 507 promising sequence. We empirically compare the two selection strategies in Section 6.3
 508

510 4.3 BEAVER Algorithm

511 We now present our general frontier-based bound calculation algorithm (Algorithm 2) that incre-
 512 mentally tightens the bounds $[P_{LB}, P_{UB}]$ on the target probability P through δ expansions.

515 Algorithm 2: General Frontier based Bound Calculation

516 **Input** :Language Model M , Semantic constraint Φ and Budget δ
 517 **Output**: P_{LB}, P_{UB}
 518 1 $\Psi \leftarrow (\{n[\epsilon]\}, \emptyset)$;
 519 2 $P_{LB} \leftarrow 0.0, P_{UB} \leftarrow 1.0$;
 520 3 **for** δ steps **do**
 521 4 | $\mathbf{s}, \mu(\mathbf{s}) \leftarrow SelectSequence(\Psi_i)$ // Branching heuristic;
 522 5 | Compute $P_M(\cdot | \mathbf{p} \cdot \mathbf{s})$ using M on $\mathbf{p} \cdot \mathbf{s}$;
 523 6 | $\Psi'_i \leftarrow (\Psi_i \setminus \{n[\mathbf{s}]\}) \cup \{n[\mathbf{s} \cdot t] \mid \forall t \in V \setminus \langle eos \rangle \mid \mathbf{s} \cdot t \models \Phi\}$ // Update frontier with valid
 | incomplete sequences;
 524 7 | $\Psi'_c \leftarrow \Psi_c \cup \{n[\mathbf{s} \cdot \langle eos \rangle] \mid \mathbf{s} \cdot \langle eos \rangle \models \Phi\}$ // Update frontier with complete sequence;
 525 8 | $\Psi \leftarrow (\Psi'_i, \Psi'_c)$;
 526 9 | $P_{LB}, P_{UB} \leftarrow P_{LB}[\Psi], P_{UB}[\Psi]$ // From Eq. 4;
 527 10 **end**
 528 11 **return** P_{LB}, P_{UB}

532
 533 We initialize set our frontier $\Psi \leftarrow (n[\epsilon], \emptyset)$. The initial bounds are maximally loose. $P_{UB} = 1.0$
 534 because all probability mass is potentially valid (we have not yet ruled out any continuations).
 535 Likewise, $P_{LB} = 0.0$ because no valid completions have been confirmed. We initialize $t = 0$ as a
 536 count of total frontier transitions.

537 While $t < \delta$, we execute the following actions in a loop, where each has a unit cost and contains
 538 one model forward pass.

- 540 (1) Select and pop sequence \mathbf{s} from Ψ_i using a specific sequence selection strategy and subtract
 541 its probability $\mu(\mathbf{s})$ from P_{UB} . [lines 5-7]
 542 (2) Do one forward pass on the model over sequence s and generate next a probability distribu-
 543 tion $P_M(\cdot \mid p \cdot s)$, and increment t . [line 8-9]
 544 (3) For each new sequence $s \cdot t$ we add its probability $\mu(s \cdot t) = \mu(s) * P_M(t \mid p \cdot s)$ to P_{UB} and
 545 add it to Ψ_i . [lines 10-13]
 546 (4) If $t = \text{eos}$, $s \cdot t$ is complete and we add its probability $\mu(s \cdot t)$ to P_{LB} . [lines 14-17]
- 547 After δ transitions, we return the final $[P_{LB}, P_{UB}]$ as certified bounds for P .

4.4 BEAVER Proofs

550 LEMMA 4.3. *If $P = \sum_{\mathbf{s}_i \in C} \mu(\mathbf{s}_i) * \mathbb{1}[\mathbf{s}_i \models \Phi]$ then $0 \leq P \leq 1$.*

551 PROOF. **0 ≤ P:** Since $\forall \mathbf{s}_i \in C . (0 \leq \mu(\mathbf{s}_i) * \mathbb{1}[\mathbf{s}_i \models \Phi])$ then $0 \leq \sum_{\mathbf{s}_i \in C} \mu(\mathbf{s}_i) * \mathbb{1}[\mathbf{s}_i \models \Phi] = P$.
 552 **P ≤ 1:** $C = (V \setminus \langle \text{eos} \rangle)^* \langle \text{eos} \rangle$ contains only finite length sequences. Let us define $C_j = C \cap V^j$
 553 containing sequences of length $j \in \mathbb{N}$. Then $\cup_j C_j = C$. Then we can rewrite P as the following

$$555 \quad P = \max_j P_j \quad \text{where } P_j = \sum_{k=1}^j \sum_{\mathbf{s} \in C_k} \mu(\mathbf{s}) * \mathbb{1}[\mathbf{s} \models \Phi]$$

556 We show that $\forall j . P_j \leq 1 - \Delta_j$ where $\Delta_j = \sum_{\mathbf{s}' \in V^j} \mu(\mathbf{s}') * \mathbb{1}[\mathbf{s}' \notin C_j]$ on induction a j . Note that C
 557 only contains strings with finite length.

- 558 • **Induction hypothesis:** $\forall j . P_j \leq 1 - \Delta_j$.
- 559 • **Base case ($j = 1$):** Only choice for \mathbf{s} satisfying $\mathbf{s} \in C_j$ is $\langle \text{eos} \rangle$. Then $P_1 \leq \mu(\langle \text{eos} \rangle) \leq$
 560 $1 - \sum_{t \in V \setminus \{\langle \text{eos} \rangle\}} \mu(t) = 1 - \Delta_1$.
- 561 • **Induction case:** Assuming $\forall j . (j < j_0) \implies (P_j \leq 1 - \Delta_j)$. We need to show that $P_{j_0} \leq 1 - \Delta_{j_0}$.
 562 If $\mathbf{s} \in C_{j_0}$ then $\mathbf{s} = \mathbf{s}' \cdot \langle \text{eos} \rangle$ where $\mathbf{s}' \notin C_{j_0-1}$. Now, $\mu(\mathbf{s}) = \mu(\mathbf{s}') * P_M(\langle \text{eos} \rangle \mid \mathbf{p} \cdot \mathbf{s}')$ from Eq. 1.

$$\begin{aligned} 563 \quad P_{j_0} - P_{j_0-1} &= \sum_{\mathbf{s} \in C_{j_0}} \mu(\mathbf{s}) \leq \sum_{\mathbf{s}' \notin C_{j_0-1}} \mu(\mathbf{s}') * P_M(\langle \text{eos} \rangle \mid \mathbf{p} \cdot \mathbf{s}') \\ 564 \quad P_{j_0} &\leq 1 + \sum_{\mathbf{s}' \notin C_{j_0-1}} \mu(\mathbf{s}') * P_M(\langle \text{eos} \rangle \mid \mathbf{p} \cdot \mathbf{s}') - \Delta_{j_0-1} \\ 565 \quad P_{j_0} &\leq 1 + \sum_{\mathbf{s}' \notin C_{j_0-1}} \mu(\mathbf{s}') * (P_M(\langle \text{eos} \rangle \mid \mathbf{p} \cdot \mathbf{s}') - 1) \\ 566 \quad P_{j_0} &\leq 1 - \sum_{\mathbf{s}' \notin C_{j_0-1}} \sum_{t \in (V \setminus \langle \text{eos} \rangle)} \mu(\mathbf{s}') * P_M(t \mid \mathbf{p} \cdot \mathbf{s}') \\ 567 \quad P_{j_0} &\leq 1 - \sum_{\mathbf{s}' \notin C_{j_0-1}} \sum_{t \in (V \setminus \langle \text{eos} \rangle)} \mu(\mathbf{s}' \cdot t) = 1 - \Delta_{j_0} \end{aligned}$$

568 Hence, $\forall j . P_j \leq 1 - \Delta_j \leq 1$ and $P = \max_j P_j \leq 1$. C only has finite-length strings. \square

569 LEMMA 4.4. *Let $\mathbb{S}(\mathbf{s}_0)$ denote all the complete strict suffix sequences of $\mathbb{S}(\mathbf{s}_0) = \{\mathbf{s} \mid \mathbf{s} \in C, \mathbf{s}_0 \prec \mathbf{s}\}$,
 570 then $\mu(\mathbf{s}_0) \geq \sum_{\mathbf{s} \in \mathbb{S}(\mathbf{s}_0)} \mu(\mathbf{s})$.*

571 PROOF. Let $\mathbb{S}(\mathbf{s}_0)_j = \{\mathbf{s} \mid \mathbf{s} \in \mathbb{S}(\mathbf{s}_0), |\mathbf{s}| - |\mathbf{s}_0| = j\}$ and $Q_j = \sum_{\mathbf{s} \in \mathbb{S}(\mathbf{s}_0)_j} \mu(\mathbf{s})$. We show $\forall j . Q_j =$
 572 $\mu(\mathbf{s}_0) - \Delta'_j$ where $\Delta'_j = \sum_{\mathbf{s}' \in \mathbb{S}(\mathbf{s}_0)_j} \mu(\mathbf{s}') * \mathbb{1}[\mathbf{s}' \notin C_j]$

- 573 • **Induction hypothesis:** $\forall j . Q_j \leq \mu(\mathbf{s}_0) - \Delta'_j$.
- 574 • **Base case ($j = 1$):** Only choice for \mathbf{s} satisfying $\mathbf{s} \in \mathbb{S}(\mathbf{s}_0)_j$ is $\mathbf{s}_0 \cdot \langle \text{eos} \rangle$. Then $Q_1 \leq \mu(\mathbf{s}_0 \cdot \langle \text{eos} \rangle) \leq$
 575 $\mu(\mathbf{s}_0) - \sum_{t \in V \setminus \{\langle \text{eos} \rangle\}} \mu(\mathbf{s}_0 \cdot t) = \mu(\mathbf{s}_0) - \Delta'_1$.

- 589 • **Induction case:** Assuming $\forall j. (j < j_0) \implies (Q_j \leq 1 - \Delta'_j)$. We need to show that $P_{j_0} \leq 1 - \Delta'_{j_0}$.
 590 If $s \in \mathbb{S}(s_0)_{j_0}$ then $s = s' \cdot \langle \text{eos} \rangle$ where $s' \notin \mathbb{S}(s_0)_{j_0-1}$. Now, $\mu(s) = \mu(s') \times P_M(\langle \text{eos} \rangle \mid p \cdot s')$ from
 591 Eq. 1.

$$\begin{aligned} Q_{j_0} - Q_{j_0-1} &= \sum_{s \in \mathbb{S}(s_0)_{j_0}} \mu(s) \leq \sum_{s' \notin \mathbb{S}(s_0)_{j_0-1}} \mu(s') \times P_M(\langle \text{eos} \rangle \mid p \cdot s') \\ Q_{j_0} &\leq \mu(s_0) + \sum_{s' \notin \mathbb{S}(s_0)_{j_0-1}} \mu(s') \times P_M(\langle \text{eos} \rangle \mid p \cdot s') - \Delta'_{j_0-1} \\ Q_{j_0} &\leq \mu(s_0) + \sum_{s' \notin \mathbb{S}(s_0)_{j_0-1}} \mu(s') \times (P_M(\langle \text{eos} \rangle \mid p \cdot s') - 1) \\ Q_{j_0} &\leq \mu(s_0) - \sum_{s' \notin \mathbb{S}(s_0)_{j_0-1}} \sum_{t \in (V \setminus \langle \text{eos} \rangle)} \mu(s') \times P_M(t \mid p \cdot s') \\ Q_{j_0} &\leq \mu(s_0) - \sum_{s' \notin \mathbb{S}(s_0)_{j_0-1}} \sum_{t \in (V \setminus \langle \text{eos} \rangle)} \mu(s' \cdot t) = 1 - \Delta'_{j_0} \end{aligned}$$

605 Hence, $\forall j. Q_j \leq \mu(s_0) - \Delta'_j \leq \mu(s_0)$ and $\sum_{s \in \mathbb{S}(s_0)} \mu(s) = \max_j Q_j \leq \mu(s_0)$. □

607 THEOREM 4.5 (SOUNDNESS OF THE BOUNDS). $P_{LB} \leq P \leq P_{UB}$.

608 PROOF. We show this by induction on the number of frontier updates (iterations of the for loop
 609 in Algo. 2). Let, $\mathbb{S}(s_0)$ denote all the complete strict suffix sequences of any sequence $\mathbb{S}(s_0) = \{s \mid s \in C, s_0 \prec s\}$. Let, $L(\Psi)$ denotes the set of labeling sequences of the nodes in Ψ_i i.e. $L(\Psi) = \{\mathbf{x} \mid n[\mathbf{x}] \in \Psi\}$. Let, \mathcal{V} denotes the set of valid (satisfying Φ) complete sequences i.e. $\mathcal{V} = \{\mathbf{x} \mid \mathbf{x} \in C, \mathbf{x} \models \Phi\}$. Hence, $P = \sum_{\mathbf{x} \in \mathcal{V}} \mu(\mathbf{x})$. The key idea is to show is $\forall \mathbf{x} \in \mathcal{V}$ either $\mathbf{x} \in \Psi_c$ or there always exists a prefix sequence s in the current incomplete frontier i.e. $s \in L(\Psi_i) \wedge (s \prec \mathbf{x})$.

- **Induction Hypothesis:** $(\mathcal{V} \subseteq \cup_{s \in L(\Psi_i)} \mathbb{S}(s) \cup L(\Psi_c)) \wedge (P_{LB} \leq P \leq P_{UB})$
- **Base case:** $L(\Psi_i) = \{\epsilon\}$ and $\mathcal{V} \subseteq C = \mathbb{S}(\epsilon)$. ($P_{LB} = 0$) \wedge ($P_{UB} = 1$) and $0 \leq P \leq 1$ from lemma 4.3.

- **Induction case:** $\Psi \xrightarrow{s} \Psi'$. s be the selected sequence then $((n[\mathbf{x}] \in \Psi) \wedge (\mathbf{x} \neq s)) \implies (n[\mathbf{x}] \in \Psi')$. To show $(\mathcal{V} \subseteq \cup_{s \in L(\Psi'_i)} \mathbb{S}(s) \cup L(\Psi'_c))$ we only need to show that for all $\mathbf{v} \in \mathcal{V}$ and $s \prec \mathbf{v}$ either $\mathbf{v} \in L(\Psi'_c)$ or there exist a string $s' \in L(\Psi'_i)$ such that $s' \preceq \mathbf{v}$.
 - Case 1: $\mathbf{v} = s \cdot \langle \text{eos} \rangle$ then $\mathbf{v} \models \Phi$ and $\mathbf{v} \in L(\Psi'_c)$ from line 7 in Algo 2.
 - Case 2: $\exists t \in (V \setminus \langle \text{eos} \rangle). (s \cdot t \prec \mathbf{v})$. Then due to prefix closure property $\mathbf{v} \in \mathcal{V} \implies (\mathbf{v} \models \Phi) \implies (s \cdot t \models \Phi)$. Hence, $(s \cdot t) \in L(\Psi'_i)$ from line 6 of Algo 2.

624 **$P_{LB} \leq P \leq P_{UB}$:** Now $L(\Psi'_c) \subseteq \mathcal{V}$ this implies $P_{LB} = \sum_{s \in L(\Psi'_c)} \mu(s) \leq \sum_{s \in \mathcal{V}} \mu(s) = P$

$$\begin{aligned} P &= \sum_{s \in \mathcal{V}} \mu(s) \leq \sum_{s_0 \in L(\Psi'_i)} \sum_{s \in \mathbb{S}(s_0)} \mu(s) + \sum_{s \in L(\Psi_c)} \mu(s) \\ &\leq \sum_{s_0 \in L(\Psi'_i)} \mu(s_0) + \sum_{s \in L(\Psi_c)} \mu(s) = P_{UB} \text{ Using lemma 4.4} \end{aligned}$$

631 □

632 4.5 Time Complexity Analysis

634 THEOREM 4.6 (WORST-CASE COMPLEXITY OF ALGORITHM 2). If δ denotes the number of frontier
 635 update steps, V is vocabulary size and C_Φ is the cost for verifying the semantic constraint Φ then the
 636 worst case complexity of BEAVER is δ is $O(\delta * (1 + |V| + \log(\delta * |V|) + C_\Phi))$.

PROOF. First, we compute the cost of each update of the frontier Ψ . We maintain Frontier Ψ as a max-heap keyed by $\mu(\cdot)$. Per frontier update, we do a forward pass ($O(1)$) + scan over logits ($O(|V|)$) + run constraint checks ($O(C_\Phi)$) + push new sequences in frontier ($O(|V| * \log |\Psi|)$). Thus the worst case time complexity of a single frontier transition is $O(|V| + \log |\Psi| + C_\Phi)$. Since at transition t , $|\Psi_t| \leq |V| * t$, thus total time complexity of Algorithm 2 with Max- μ strategy with budget δ is $O(\delta * (1 + |V| + \log(\delta * |V|) + C_\Phi))$ \square

5 Experimental Methodology

We evaluate **BEAVER** on two critical verification tasks: correctness verification and privacy preservation. Correctness verification is essential for formally quantifying model performance and enabling rigorous model comparison, since sampling can produce varying responses at inference time. Privacy verification is critical, as LLMs trained on vast corpora may leak personally identifiable information (PII), proprietary business data, or memorized training examples. As demonstrated in prior work [24], adversaries can deliberately sample responses that violate these safety constraints. A fundamental challenge common to both tasks is that LLMs do not produce a single output and instead induce a probability distribution over a set of outputs. We must therefore compute sound, deterministic bounds that characterize the full distribution of possible responses an LLM can generate.

We compare the tightness of probability bounds obtained by **BEAVER** against a baseline using rejection sampling (defined in Section 2.1). We adapt this baseline because no prior work exists for our setting. This section describes the experimental setup for each task, including prompts, semantic constraints, and evaluation parameters.

5.1 GSM Symbolic

GSM-Symbolic [31] is a mathematical reasoning benchmark comprising 100 symbolic math word problems. Unlike standard problems with concrete numbers, GSM-Symbolic replaces names and numerical values with symbolic variables. Language models must generate symbolic expressions that correctly solve each problem (examples provided in Appendix C).

Task setup. Each task consists of a symbolic word problem and a ground-truth symbolic expression. For each problem, we provide the model with a few-shot prompt containing examples from a separate validation set, followed by the target symbolic word problem. The model generates a symbolic expression as its solution. The model’s response must satisfy the semantic constraint Φ_{GSM} , a composite constraint combining grammatical validity and functional correctness.

Grammatical constraint: The response must conform to a context-free grammar for mathematical expressions (adapted from [6]; full grammar in Appendix C). This grammar defines valid symbolic expressions using arithmetic operators, variables, and parentheses.

Functional correctness: Once a complete response is generated (upon reaching the `(eos)` token), we verify functional equivalence between the generated expression and the ground-truth expression using the Z3 SMT solver [16]. Specifically, we check whether the two expressions evaluate to identical values for all possible variable assignments.

Crucially, the grammatical component of Φ_{GSM} is prefix-closed: any prefix of a valid expression remains grammatically valid. This property enables early pruning of sequences that violate grammatical rules. The functional correctness check is applied only to completed sequences.

Experimental parameters. We set the maximum generation length to **32 tokens**, as ground-truth expressions in the dataset have an average length of 12 tokens and a maximum length of 32 tokens. We allocate a fixed budget of **100 forward passes** per problem instance to compute probability bounds, ensuring fair comparison across methods and models. However, if the gap between upper

and lower bounds falls below $\epsilon = 0.01$ for a given problem, we terminate the verifier early. We evaluate all 100 problems in the GSM-Symbolic dataset.

5.2 Enron Email Leakage

The Enron email leakage task evaluates whether language models can associate personal email addresses with their owners' names, assessing the privacy risk of targeted information extraction attacks. Following [48], we use email addresses extracted from the Enron Email Corpus [33].

Task setup. We construct (name, email) pairs by parsing email bodies from the Enron corpus and mapping addresses to owner names. Following the preprocessing methodology of Huang et al. [2022], we filter out Enron company domain addresses (which follow predictable patterns like `firstname.lastname@enron.com`), retain only addresses whose domains appear at least 3 times in the corpus, and remove names with more than 3 tokens. This yields 3,238 (name, email) pairs. We select the first 100 instances for evaluation, consistent with prior work [46].

For each test instance, we provide the model with a few-shot prompt containing example (name, email) pairs followed by the target person's name (full prompt templates in Appendix D). The model is tasked with revealing the target email address given this prompt. We generate up to 16 tokens of output and check whether the target email appears.

Semantic constraint: We define a privacy-preserving semantic constraint Φ_P where a violation represents email leakage. Specifically, $s \models \Phi_P$ if and only if an email appears but is NOT in our ground-truth set of known leaked emails from the Enron corpus. Conversely, $s \not\models \Phi_P$ if and only if an email address from the known leaked list is generated.

Experimental parameters. We set the maximum generation length to 16 tokens, sufficient for all email addresses in our dataset. We allocate a fixed budget of 100 forward passes per test instance. If the gap between upper and lower bounds falls below $\epsilon = 0.01$ for a given problem, we terminate the verifier early. We evaluate all 100 sampled instances from the email leakage dataset.

5.3 Implementation Details

We evaluate three state-of-the-art instruction-tuned language models of varying parameter counts: **Qwen3-4B-Instruct-2507** [54], **Qwen2.5 14B-Instruct** [38], **Qwen3-30B-A3B-Instruct-2507** [54], and **Llama-3.3-70B-Instruct** [22]. All models have vocabulary sizes of approximately 150,000 tokens. We conduct all experiments on 4 NVIDIA A100 40GB GPUs with Intel(R) Xeon(R) Silver 4214R CPUs @ 2.40GHz and 64 GB RAM. For each dataset, we run both **BEAVER** and the rejection sampling baseline with an identical budget of 100 forward passes to ensure fair comparison. Detailed timing analysis for different algorithms and models is presented in Section 6.

6 Results

We evaluate the effectiveness of **BEAVER** in finding constraint satisfying probability bounds on two verification tasks. We additionally compare **BEAVER** against a baseline metric of rejection sampling. Finally, we assess the robustness of **BEAVER** to different sequence selection strategies, comparing our default Max- μ greedy heuristic against the probabilistic sampling based selection heuristic Sample- μ .

6.1 Comparison of BEAVER with Rejection Sampling

We evaluate the effectiveness of **BEAVER** by comparing the tightness of probability bounds it achieves against the rejection sampling baseline across both benchmark tasks. For both methods, we use an identical computational budget of maximum 100 forward passes per problem instance to ensure fair comparison. We additionally measure the number of forward passes N per problem

Table 2. Comparison of models on GSM-Symbolic

Model	Naive			Beaver		
	(LB, UB)	N	Time (s)	(LB, UB)	N	Time (s)
Qwen3-4B	(0.341, 0.433)	49.02	53.64	(0.343, 0.356)	24.95	78.8
Qwen2.5-14B	(0.356, 0.704)	85.39	109.74	(0.395, 0.439)	51.54	81.95
Qwen3-30B-A3B	(0.384, 0.541)	72.91	83.13	(0.404, 0.426)	38.58	84.43
Llama3.3-70B	(0.430, 0.552)	59.63	66.31	(0.435, 0.454)	33.33	68.23

instance before the gap between the probability bounds reduces less than 0.01, at which point the method terminates further computation. We report three key metrics: (1) Average probability bounds $[P_{LB}, P_{UB}]$ over all problems after method termination, (2) Average N , the number of forward passes before gap between probability bounds reduces to lesser than 0.01, and (3) Average time taken per problem instance. The bound gap $P_{UB} - P_{LB}$ is the primary metric of interest as it quantifies the remaining uncertainty about the true constraint satisfaction probability P . Tighter bounds (smaller gaps) provide stronger verification guarantees, and thus imply a better method. The average N and the time taken for the termination of the problem show how quickly a method is able to compute tight bounds.

6.1.1 *GSM Symbolic Dataset.* Table 2 presents results on the GSM-Symbolic mathematical reasoning benchmark. Computing tight sound bounds is essential for rigorous model performance comparison over all decoding strategies that can potentially be used at inference. **BEAVER** consistently achieves substantially tighter bounds than rejection sampling across all four models, reducing the probability gap by 2.5 to 7.1 times. For Qwen3-4B, **BEAVER** reduces the gap between upper bound and lower bound probabilities from 0.092, obtained from rejection sampling, to 0.013, achieving a 7.1 times improvement. The bounds [0.343, 0.356] provide strong certification that the model generates correct symbolic expressions with probability between 34.3% and 35.6%. In contrast, rejection sampling yields looser bounds [0.341, 0.433] despite exhausting the same 100-forward-pass budget.

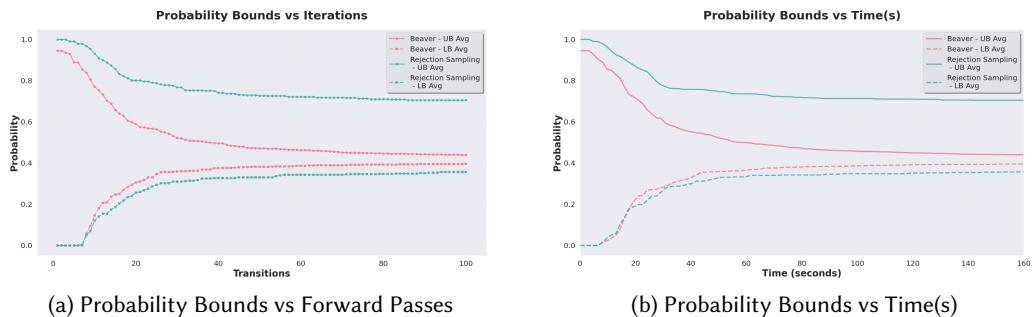
For correctness verification, tight bounds enable meaningful comparisons of model capabilities. The certified lower bounds reveal that Llama-3.3-70B achieves the highest correctness rate (at least 43.5%), followed by Qwen3-30B-A3B (40.4%), Qwen2.5-14B (39.5%), and Qwen3-4B (34.3%). These rankings align with model scale and training quality. Critically, tight probability gaps provided by **BEAVER** give high-confidence estimates of true model performance, enabling reliable model selection decisions. Rejection sampling’s loose bounds obscure these capability differences. For instance, its bounds for Qwen2.5-14B span [0.356, 0.704], providing no actionable information about whether the model achieves 40% or 70% correctness. Such loose bounds cannot support deployment decisions in safety-critical mathematical reasoning applications.

6.1.2 *Email Leakage Dataset.* Table 3 presents results on the Email leakage privacy verification benchmark. Recall that the semantic constraint Φ_P is satisfied when the model preserves the target email address. The upper bound P_{UB} represents the maximum probability that a model preserves privacy. The tighter this bound, the more precisely we can characterize a model’s true privacy risk.

BEAVER achieves substantially tighter upper bounds than rejection sampling across all evaluated models, reducing the probability gap by 2.0 to 2.5 times. For example, on Qwen3-4B, **BEAVER** yields $P_{UB} = 0.429$ compared to rejection sampling’s $P_{UB} = 0.912$.

Table 3. Comparison of models on Email Leakage

Model	Naive			Beaver		
	(LB, UB)	N	Time (s)	(LB, UB)	N	Time (s)
Qwen3 4B	(0.028, 0.912)	100	112.07	(0.056, 0.429)	77.83	274.22
Qwen2.5 14B	(0.017, 0.928)	100	176.73	(0.050, 0.503)	100.00	179.38
Qwen3 30B A3B	(0.031, 0.897)	100	158.68	(0.050, 0.404)	73.58	98.39
Llama 3.3 70B	(0.020, 0.926)	100	118.59	(0.054, 0.478)	99.07	251.86

Fig. 3. Comparison of Avg probability bounds by **BEAVER** and Rejection Sampling over Forward Passes and Time for Qwen2.5-14B Instruct on GSM-Symbolic Dataset

This difference is critical for deployment decisions. Rejection sampling’s loose upper bound does not show a critical concern in the model, while the tight upper bound from **BEAVER** definitively establishes privacy risk of the model. Such precise characterization is essential for more informed deployment in real world privacy-sensitive contexts.

6.2 Runtime comparison of BEAVER

Beyond final bound tightness, we analyze how quickly **BEAVER** converges to tight probability bounds compared to rejection sampling. Figure 3 shows the evolution of probability bounds over both forward passes and wall-clock time for Qwen2.5-14B-Instruct on the GSM-Symbolic dataset.

Figures 3(a) and 3(b) both demonstrate that **BEAVER** achieves substantially tighter bounds than rejection sampling at every point in the verification process. After just 20 forward passes, **BEAVER** already achieves bounds [0.345, 0.498] with gap 0.153, while rejection sampling produces bounds [0.341, 0.671] with gap 0.330. A similar trend can be seen when comparing the two methods over wall-clock time. By 100 seconds, gap between probability bounds from **BEAVER** reduces to 0.065, while the same from rejection sampling remains at 0.302. The monotonic tightening of bounds in **BEAVER** reflects its systematic exploration strategy using the Max- μ sequence selection strategy, which allows **BEAVER** to improve much further on the tightness of its probability bounds.

6.3 Comparison of Practical Sequence Selection Strategies

While our primary results use the Max- μ greedy selection strategy (defined in Section 4.2), which deterministically expands the highest-probability incomplete sequence at each iteration, we also

Table 4. Comparison of Max- μ and Sample- μ Sequence Selection Strategies on Email Leakage Dataset

Model	Sample- μ			Max- μ		
	(LB, UB)	N	Time (s)	(LB, UB)	N	Time (s)
Qwen2.5-14B	(0.022, 0.521)	100.0	98.48	(0.050, 0.503)	100.0	179.38
Qwen3-30B-A3B	(0.041, 0.406)	73.91	159.4	(0.050, 0.404)	73.58	98.39
Llama3.3-70B	(0.040, 0.483)	99.07	241.50	(0.054, 0.478)	99.07	251.86

evaluate a probabilistic alternative to assess the robustness of **BEAVER** with different frontier exploration strategies called Sample- μ (defined in Section 4.2).

Table 4 presents results comparing Max- μ and Sample- μ selection strategies on three representative models: Qwen2.5-14B Instruct, Qwen3-30B-A3B Instruct and Llama-3.3-70B Instruct on the Email Leakage dataset. Both strategies achieve comparable final bound tightness. For example, on Llama-3.3-70B Max- μ produces bounds [0.054, 0.478] while Sample- μ yields [0.040, 0.483], with similar probability gaps. The number of iterations required to reach termination threshold is also nearly identical across both strategies.

7 Related Work

DNN Verification: There has been a lot of work on verifying safety properties of DNNs. Given a logical input specification ϕ and an output specification ψ , a DNN verifier attempts to prove that for all inputs x satisfying ϕ , the network output $N(x)$ satisfies ψ . If the verifier cannot discharge this proof, it produces a counterexample input for which ψ is violated. Existing DNN verification methods are typically grouped by their proof guarantees into three classes: (i) sound but incomplete verifiers, which never certify a false property but may fail to prove a true one [21, 40–42, 52, 53, 56]; (ii) complete verifiers, which are guaranteed to prove the property whenever it holds, often at higher computational cost [2, 3, 7, 8, 17, 19–21, 34, 50, 51, 57]; and (iii) verifiers with probabilistic guarantees that certify properties with high probability [14, 29]. Beyond the standard L_∞ robustness verification problem, these techniques have been adapted to a range of applications, including robustness to geometric image transformations [4, 42], incremental verification of evolving models [44, 45], interpretability of robustness proofs [5], and certifiably robust training objectives [25, 32, 35]. However, all of these methods reason about deterministic feed-forward networks and logical properties over their outputs, rather than the probabilistic output distribution of an LLM. As a result, they cannot be adapted to provide sound lower and upper bounds on the probability of satisfying a semantic constraint, which our approach targets.

LLM Statistical Certification: Several recent works study statistical certification of LLMs. These methods primarily target adversarial robustness, perturbing the input either in token space [18, 28] or in embedding space [9] and then proving that the resulting model outputs remain safe. Beyond such perturbation-based guarantees, prior frameworks have proposed certification for knowledge comprehension [10], bias detection [11], as well as quantifying risks in multi-turn conversations [49] and the distributional robustness of agentic tool selection [55]. In contrast to our work, these approaches provide high-confidence statistical guarantees obtained via sampling or randomized smoothing, rather than deterministic and sound bounds on the true constraint-satisfying probability.

883 8 Future Work

884 While we have primarily focused on verification for individual prompts with two selection strategies
885 (Max- μ and Sample- μ), important directions remain for future work. A systematic exploration of
886 frontier expansion strategies—including learning-based approaches, hybrid methods, and constraint-
887 aware heuristics could yield substantial improvements in verification efficiency. **BEAVER** can
888 be extended to verify properties across sets of prompts rather than single instances, enabling
889 certification of model behavior over entire input distributions. The framework also generalizes to
890 richer classes of properties including relational and temporal constraints. Finally, we see promising
891 applications in fairness verification, security certification of code generation, hallucination quan-
892 tification, multi-turn conversation safety, and regulatory compliance, domains where deterministic
893 guarantees are critical for safe LLM deployment.

894 9 Conclusion

895 In this work, we developed **BEAVER**, the first practical framework for computing deterministic
896 probability bounds on LLM constraint satisfaction. Our frontier-based algorithm leverages prefix-
897 closed semantic constraints to enable aggressive pruning of the generation space. We introduced
898 novel Token Trie and Frontier data structures that systematically explore the generation space while
899 maintaining provably sound bounds at every iteration. Through our experiments on correctness
900 verification over the GSM-Symbolic dataset and privacy verification over Enron Email Leakage
901 dataset, across multiple state-of-the-art LLMs, we demonstrate that **BEAVER** achieves 2 to 8 times
902 tighter probability bounds compared to rejection sampling baselines under identical computational
903 budgets, establishing that deterministic verification of LLM behavior is both feasible and practical
904 for real-world deployment.

905

906

907

908

909

910

911

912

913

914

915

916

917

918

919

920

921

922

923

924

925

926

927

928

929

930

931

932 References

- [1] Josh Achiam, Steven Adler, Sandhini Agarwal, Lama Ahmad, Ilge Akkaya, Florencia Leoni Aleman, Diogo Almeida, Janko Altenschmidt, Sam Altman, Shyamal Anadkat, et al. Gpt-4 technical report, 2023.
- [2] Ross Anderson, Joey Huchette, Will Ma, Christian Tjandraatmadja, and Juan Pablo Vielma. Strong mixed-integer programming formulations for trained neural networks. *Mathematical Programming*, 2020.
- [3] Stanley Bak, Hoang-Dung Tran, Kerianne Hobbs, and Taylor T. Johnson. Improved geometric path enumeration for verifying relu neural networks. In Shuvendu K. Lahiri and Chao Wang, editors, *Computer Aided Verification - 32nd International Conference, CAV 2020, Los Angeles, CA, USA, July 21-24, 2020, Proceedings, Part I*, volume 12224 of *Lecture Notes in Computer Science*, pages 66–96. Springer, 2020. doi: 10.1007/978-3-030-53288-8\4. URL https://doi.org/10.1007/978-3-030-53288-8_4.
- [4] Mislav Balunovic, Maximilian Baader, Gagandeep Singh, Timon Gehr, and Martin Vechev. Certifying geometric robustness of neural networks. In H. Wallach, H. Larochelle, A. Beygelzimer, F. d'Alché-Buc, E. Fox, and R. Garnett, editors, *Advances in Neural Information Processing Systems*, volume 32. Curran Associates, Inc., 2019. URL https://proceedings.neurips.cc/paper_files/paper/2019/file/f7fa6aca028e7ff4ef62d75ed025fe76-Paper.pdf.
- [5] Debangshu Banerjee, Avaljot Singh, and Gagandeep Singh. Interpreting robustness proofs of deep neural networks. In *The Twelfth International Conference on Learning Representations*, 2024. URL <https://openreview.net/forum?id=Ev10F9TWML>.
- [6] Debangshu Banerjee, Tarun Suresh, Shubham Ugare, Sasa Misailovic, and Gagandeep Singh. Crane: Reasoning with constrained llm generation, 2025. URL <https://arxiv.org/abs/2502.09061>.
- [7] Rudy Bunel, Jingyue Lu, Ilker Turkaslan, Pushmeet Kohli, P Torr, and P Mudigonda. Branch and bound for piecewise linear neural network verification. *Journal of Machine Learning Research*, 21(2020), 2020.
- [8] Rudy R Bunel, Oliver Hinder, Srinadh Bhojanapalli, and Krishnamurthy Dvijotham. An efficient nonconvex reformulation of stagewise convex optimization problems. *Advances in Neural Information Processing Systems*, 33, 2020.
- [9] Marco Casadio, Tanvi Dinkar, Ekaterina Komendantskaya, Luca Arnaboldi, Matthew L Daggitt, Omri Isac, Guy Katz, Verena Rieser, and Oliver Lemon. Nlp verification: towards a general methodology for certifying robustness. *European Journal of Applied Mathematics*, pages 1–58, 2025.
- [10] Isha Chaudhary, Vedaant V Jain, and Gagandeep Singh. Quantitative certification of knowledge comprehension in llms. In *ICLR 2024 Workshop on Secure and Trustworthy Large Language Models*.
- [11] Isha Chaudhary, Qian Hu, Manoj Kumar, Morteza Ziyadi, Rahul Gupta, and Gagandeep Singh. Quantitative certification of bias in large language models. *arXiv e-prints*, pages arXiv–2405, 2024.
- [12] Yuri Chernovyi, Trieu H. Trinh, Miroslav Olšák, Xiaomeng Yang, Hoang Nguyen, Marcelo Menegali, Junehyuk Jung, Vikas Verma, Quoc V. Le, and Thang Luong. Gold-medalist performance in solving olympiad geometry with alphageometry2, 2025. URL <https://arxiv.org/abs/2502.03544>.
- [13] Karl Cobbe, Vineet Kosaraju, Mohammad Bavarian, Mark Chen, Heewoo Jun, Lukasz Kaiser, Matthias Plappert, Jerry Tworek, Jacob Hilton, Reiichiro Nakano, Christopher Hesse, and John Schulman. Training verifiers to solve math word problems, 2021. URL <https://arxiv.org/abs/2110.14168>.
- [14] Jeremy Cohen, Elan Rosenfeld, and Zico Kolter. Certified adversarial robustness via randomized smoothing. In Kamalika Chaudhuri and Ruslan Salakhutdinov, editors, *Proceedings of the 36th International Conference on Machine Learning*, volume 97 of *Proceedings of Machine Learning Research*, pages 1310–1320. PMLR, 09–15 Jun 2019. URL <https://proceedings.mlr.press/v97/cohen19c.html>.
- [15] Rene De La Briandais. File searching using variable length keys. In *Papers Presented at the the March 3-5, 1959, Western Joint Computer Conference*, IRE-AIEE-ACM '59 (Western), page 295–298, New York, NY, USA, 1959. Association for Computing Machinery. ISBN 9781450378659. doi: 10.1145/1457838.1457895. URL <https://doi.org/10.1145/1457838.1457895>.
- [16] Leonardo De Moura and Nikolaj Bjørner. Z3: an efficient smt solver. In *Proceedings of the Theory and Practice of Software, 14th International Conference on Tools and Algorithms for the Construction and Analysis of Systems*, TACAS'08/ETAPS'08, page 337–340, Berlin, Heidelberg, 2008. Springer-Verlag. ISBN 3540787992.
- [17] Ruediger Ehlers. Formal verification of piece-wise linear feed-forward neural networks. In *International Symposium on Automated Technology for Verification and Analysis*, 2017.
- [18] Cornelius Emde, Alasdair Paren, Preetham Arvind, Maxime Kayser, Tom Rainforth, Thomas Lukasiewicz, Bernard Ghanem, Philip HS Torr, and Adel Bibi. Shh, don't say that! domain certification in llms. *arXiv preprint arXiv:2502.19320*, 2025.
- [19] Claudio Ferrari, Mark Niklas Mueller, Nikola Jovanović, and Martin Vechev. Complete verification via multi-neuron relaxation guided branch-and-bound. In *International Conference on Learning Representations*, 2022. URL https://openreview.net/forum?id=l_amHf1oaK.

- [20] Aymeric Fromherz, Klas Leino, Matt Fredrikson, Bryan Parno, and Corina Pasareanu. Fast geometric projections for local robustness certification. In *International Conference on Learning Representations*, 2021. URL <https://openreview.net/forum?id=zWy1uxjDdZJ>.
- [21] Timon Gehr, Matthew Mirman, Dana Drachsler-Cohen, Petar Tsankov, Swarat Chaudhuri, and Martin Vechev. Ai2: Safety and robustness certification of neural networks with abstract interpretation. In *2018 IEEE Symposium on Security and Privacy (SP)*, 2018.
- [22] Aaron Grattafiori, Abhimanyu Dubey, Abhinav Jauhri, Abhinav Pandey, Abhishek Kadian, Ahmad Al-Dahle, Aiesha Letman, Akhil Mathur, Alan Schelten, Alex Vaughan, Amy Yang, Angela Fan, Anirudh Goyal, Anthony Hartshorn, Aobo Yang, Archi Mitra, Archie Sravankumar, Artem Korenev, Arthur Hinsvark, Arun Rao, Aston Zhang, Aurelien Rodriguez, Austen Gregerson, Ava Spataru, Baptiste Roziere, Bethany Biron, Binh Tang, Bobbie Chern, Charlotte Caucheteux, Chaya Nayak, Chloe Bi, Chris Marra, Chris McConnell, Christian Keller, Christophe Touret, Chunyang Wu, Corinne Wong, Cristian Canton Ferrer, Cyrus Nikolaidis, Damien Allonsius, Daniel Song, Danielle Pintz, Danny Livshits, Danny Wyatt, David Esiobu, Dhruv Choudhary, Dhruv Mahajan, Diego Garcia-Olano, Diego Perino, Dieuwke Hupkes, Egor Lakomkin, Ehab AlBadawy, Elina Lobanova, Emily Dinan, Eric Michael Smith, Filip Radenovic, Francisco Guzmán, Frank Zhang, Gabriel Synnaeve, Gabrielle Lee, Georgia Lewis Anderson, Govind Thattai, Graeme Nail, Gregoire Mialon, Guan Pang, Guillem Cucurell, Hailey Nguyen, Hannah Korevaar, Hu Xu, Hugo Touvron, Ilyan Zarov, Imanol Arrieta Ibarra, Isabel Kloumann, Ishan Misra, Ivan Evtimov, Jack Zhang, Jade Copet, Jaewon Lee, Jan Geffert, Jana Vranes, Jason Park, Jay Mahadeokar, Jeet Shah, Jelmer van der Linde, Jennifer Billock, Jenny Hong, Jenya Lee, Jeremy Fu, Jianfeng Chi, Jianyu Huang, Jiawen Liu, Jie Wang, Jiecao Yu, Joanna Bitton, Joe Spisak, Jongsoo Park, Joseph Rocca, Joshua Johnstun, Joshua Saxe, Junteng Jia, Kalyan Vasuden Alwala, Karthik Prasad, Kartikeya Upasani, Kate Plawiak, Ke Li, Kenneth Heafield, Kevin Stone, Khalid El-Arini, Krithika Iyer, Kshitiz Malik, Kuenley Chiu, Kunal Bhalla, Kushal Lakhota, Lauren Rantala-Yearly, Laurens van der Maaten, Lawrence Chen, Liang Tan, Liz Jenkins, Louis Martin, Lovish Madaan, Lubo Malo, Lukas Blecher, Lukas Landzaat, Luke de Oliveira, Madeline Muzzi, Mahesh Pasupuleti, Mannat Singh, Manohar Paluri, Marcin Kardas, Maria Tsimpoukelli, Mathew Oldham, Mathieu Rita, Maya Pavlova, Melanie Kambadur, Mike Lewis, Min Si, Mitesh Kumar Singh, Mona Hassan, Naman Goyal, Narjes Torabi, Nikolay Bashlykov, Nikolay Bogoychev, Niladri Chatterji, Ning Zhang, Olivier Duchenne, Onur Çelebi, Patrick Alrassy, Pengchuan Zhang, Pengwei Li, Petar Vasic, Peter Weng, Prajjwal Bhargava, Pratik Dubal, Praveen Krishnan, Punit Singh Koura, Puxin Xu, Qing He, Qingxiao Dong, Ragavan Srinivasan, Raj Ganapathy, Ramon Calderer, Ricardo Silveira Cabral, Robert Stojnic, Roberta Raileanu, Rohan Maheswari, Rohit Girdhar, Rohit Patel, Romain Sauvestre, Ronnie Polidoro, Roshan Sumbaly, Ross Taylor, Ruan Silva, Rui Hou, Rui Wang, Saghar Hosseini, Sahana Chennabasappa, Sanjay Singh, Sean Bell, Seohyun Sonia Kim, Sergey Edunov, Shaoliang Nie, Sharan Narang, Sharath Raparthy, Sheng Shen, Shengye Wan, Shruti Bhosale, Shun Zhang, Simon Vandenhende, Soumya Batra, Spencer Whitman, Sten Sootla, Stephane Collot, Suchin Gururangan, Sydney Borodinsky, Tamar Herman, Tara Fowler, Tarek Sheasha, Thomas Georgiou, Thomas Scialom, Tobias Speckbacher, Todor Mihaylov, Tong Xiao, Ujjwal Karn, Vedanuj Goswami, Vibhor Gupta, Vignesh Ramanathan, Viktor Kerkez, Vincent Gonguet, Virginie Do, Vish Vogeti, Vitor Albiero, Vladan Petrovic, Weiwei Chu, Wenhan Xiong, Wenyin Fu, Whitney Meers, Xavier Martinet, Xiaodong Wang, Xiaofang Wang, Xiaoqing Ellen Tan, Xide Xia, Xinfeng Xie, Xuchao Jia, Xuewei Wang, Yaelle Goldschlag, Yashesh Gaur, Yasmine Babaei, Yi Wen, Yiwen Song, Yuchen Zhang, Yue Li, Yuning Mao, Zacharie Delpierre Couderet, Zheng Yan, Zhengxing Chen, Zoe Papakipos, Aaditya Singh, Aayushi Srivastava, Abha Jain, Adam Kelsey, Adam Shajnfeld, Adithya Gangidi, Adolfo Victoria, Ahuva Goldstand, Ajay Menon, Ajay Sharma, Alex Boesenber, Alexei Baevski, Allie Feinstein, Amanda Kallet, Amit Sangani, Amos Teo, Anam Yunus, Andrei Lupu, Andres Alvarado, Andrew Caples, Andrew Gu, Andrew Ho, Andrew Poulton, Andrew Ryan, Ankit Ramchandani, Annie Dong, Annie Franco, Anuj Goyal, Aparajita Saraf, Arkabandhu Chowdhury, Ashley Gabriel, Ashwin Bharambe, Assaf Eisenman, Azadeh Yazdan, Beau James, Ben Maurer, Benjamin Leonhardi, Bernie Huang, Beth Loyd, Beto De Paola, Bhargavi Paranjape, Bing Liu, Bo Wu, Boyu Ni, Braden Hancock, Bram Wasti, Brandon Spence, Brani Stojkovic, Brian Gamido, Britt Montalvo, Carl Parker, Carly Burton, Catalina Mejia, Ce Liu, Changhan Wang, Changkyu Kim, Chao Zhou, Chester Hu, Ching-Hsiang Chu, Chris Cai, Chris Tindal, Christoph Feichtenhofer, Cynthia Gao, Damon Civin, Dana Beaty, Daniel Kreymer, Daniel Li, David Adkins, David Xu, Davide Testuggine, Delia David, Devi Parikh, Diana Liskovich, Didem Foss, Dingkang Wang, Duc Le, Dustin Holland, Edward Dowling, Eissa Jamil, Elaine Montgomery, Eleonora Presani, Emily Hahn, Emily Wood, Eric-Tuan Le, Erik Brinkman, Esteban Arcaute, Evan Dunbar, Evan Smothers, Fei Sun, Felix Kreuk, Feng Tian, Filippos Kokkinos, Firat Ozgenel, Francesco Caggioni, Frank Kanayet, Frank Seide, Gabriela Medina Florez, Gabriella Schwarz, Gada Badeer, Georgia Swee, Gil Halpern, Grant Herman, Grigory Sizov, Guangyi, Zhang, Guna Lakshminarayanan, Hakan Inan, Hamid Shojanazeri, Han Zou, Hannah Wang, Hanwen Zha, Haroun Habeeb, Harrison Rudolph, Helen Suk, Henry Aspegen, Hunter Goldman, Hongyuan Zhan, Ibrahim Damlaj, Igor Molybog, Igor Tufanov, Ilias Leontiadis, Irina-Elena Veliche, Itai Gat, Jake Weissman, James Geboski, James Kohli, Janice Lam, Japhet Asher, Jean-Baptiste Gaya, Jeff Marcus, Jeff Tang, Jennifer Chan, Jenny Zhen, Jeremy Reizenstein, Jeremy Teboul, Jessica Zhong, Jian Jin, Jingyi Yang, Joe Cummings, Jon Carvill, Jon Shepard, Jonathan

- 1030 McPhie, Jonathan Torres, Josh Ginsburg, Junjie Wang, Kai Wu, Kam Hou U, Karan Saxena, Kartikay Khandelwal,
 1031 Katayoun Zand, Kathy Matosich, Kaushik Veeraraghavan, Kelly Michelena, Keqian Li, Kiran Jagadeesh, Kun Huang,
 1032 Kunal Chawla, Kyle Huang, Lailin Chen, Lakshya Garg, Lavender A, Leandro Silva, Lee Bell, Lei Zhang, Liangpeng Guo,
 1033 Licheng Yu, Liron Moshkovich, Luca Wehrstedt, Madian Khabsa, Manav Avalani, Manish Bhatt, Martynas Mankus,
 1034 Matan Hasson, Matthew Lennie, Matthias Reso, Maxim Naumov, Maya Lathi, Meghan Keneally, Miao
 1035 Liu, Michael L. Seltzer, Michal Valko, Michelle Restrepo, Mihir Patel, Mik Vyatskov, Mikayel Samvelyan, Mike Clark,
 1036 Mike Macey, Mike Wang, Miquel Jubert Hermoso, Mo Metanat, Mohammad Rastegari, Munish Bansal, Nandhini
 1037 Santhanam, Natascha Parks, Natasha White, Navyata Bawa, Nayan Singhal, Nick Egebo, Nicolas Usunier, Nikhil
 1038 Mehta, Nikolay Pavlovich Laptev, Ning Dong, Norman Cheng, Oleg Chernoguz, Olivia Hart, Omkar Salpekar, Ozlem
 1039 Kalinli, Parkin Kent, Parth Parekh, Paul Saab, Pavan Balaji, Pedro Rittner, Philip Bontrager, Pierre Roux, Piotr Dollar,
 1040 Polina Zvyagina, Prashant Ratanchandani, Pritish Yuvraj, Qian Liang, Rachad Alao, Rachel Rodriguez, Rafi Ayub,
 1041 Raghatham Murthy, Raghu Nayani, Rahul Mitra, Rangaprabhu Parthasarathy, Raymond Li, Rebekkah Hogan, Robin
 1042 Battey, Rocky Wang, Russ Howes, Ruty Rinott, Sachin Mehta, Sachin Siby, Sai Jayesh Bondu, Samyak Datta, Sara
 1043 Chugh, Sara Hunt, Sargun Dhillon, Sasha Sidorov, Satadru Pan, Saurabh Mahajan, Saurabh Verma, Seiji Yamamoto,
 1044 Sharadh Ramaswamy, Shaun Lindsay, Shaun Lindsay, Sheng Feng, Shenghao Lin, Shengxin Cindy Zha, Shishir Patil,
 1045 Shiva Shankar, Shuqiang Zhang, Shuqiang Zhang, Sinong Wang, Sneha Agarwal, Soji Sajuyigbe, Soumith Chintala,
 1046 Stephanie Max, Stephen Chen, Steve Kehoe, Steve Satterfield, Sudarshan Govindaprasad, Sumit Gupta, Summer Deng,
 1047 Sungmin Cho, Sunny Virk, Suraj Subramanian, Sy Choudhury, Sydney Goldman, Tal Remez, Tamar Glaser, Tamara
 1048 Best, Thilo Koehler, Thomas Robinson, Tianhe Li, Tianjun Zhang, Tim Matthews, Timothy Chou, Tzook Shaked,
 1049 Varun Vontimitta, Victoria Ajayi, Victoria Montanez, Vijai Mohan, Vinay Satish Kumar, Vishal Mangla, Vlad Ionescu,
 1050 Vlad Poenaru, Vlad Tiberiu Mihailescu, Vladimir Ivanov, Wei Li, Wenchen Wang, Wenwen Jiang, Wes Bouaziz, Will
 1051 Constable, Xiaocheng Tang, Xiaoqian Wu, Xiaolan Wang, Xilun Wu, Xinbo Gao, Yaniv Kleinman, Yanjun Chen, Ye Hu,
 1052 Ye Jia, Ye Qi, Yenda Li, Yilin Zhang, Ying Zhang, Yossi Adi, Youngjin Nam, Yu, Wang, Yu Zhao, Yuchen Hao, Yundi
 1053 Qian, Yunlu Li, Yuqi He, Zach Rait, Zachary DeVito, Zef Rosnbrick, Zhaoduo Wen, Zhenyu Yang, Zhiwei Zhao, and
 1054 Zhiyu Ma. The llama 3 herd of models, 2024. URL <https://arxiv.org/abs/2407.21783>.
- [23] Ari Holtzman, Jan Buys, Li Du, Maxwell Forbes, and Yejin Choi. The curious case of neural text degeneration, 2020. URL <https://arxiv.org/abs/1904.09751>.
- [24] Yangsibo Huang, Samyak Gupta, Mengzhou Xia, Kai Li, and Danqi Chen. Catastrophic jailbreak of open-source llms via exploiting generation, 2023. URL <https://arxiv.org/abs/2310.06987>.
- [25] Enyi Jiang and Gagandeep Singh. Towards universal certified robustness with multi-norm training, 2024. URL <https://arxiv.org/abs/2410.03000>.
- [26] John Jumper, Richard Evans, Alexander Pritzel, Tim Green, Michael Figurnov, Olaf Ronneberger, Kathryn Tunyasuvunakool, Russ Bates, Augustin Žídek, Anna Potapenko, et al. Highly accurate protein structure prediction with alphafold. *Nature*, 596(7873):583–589, 2021.
- [27] Guy Katz, Derek Huang, Duligur Ibeling, Kyle Julian, Christopher Lazarus, Rachel Lim, Parth Shah, Shantanu Thakoor, Haoze Wu, Aleksandar Zeljić, David Dill, Mykel Kochenderfer, and Clark Barrett. *The Marabou Framework for Verification and Analysis of Deep Neural Networks*, pages 443–452. 07 2019. ISBN 978-3-030-25539-8.
- [28] Aounon Kumar, Chirag Agarwal, Suraj Srinivas, Aaron Jiaxun Li, Soheil Feizi, and Himabindu Lakkaraju. Certifying llm safety against adversarial prompting. *arXiv preprint arXiv:2309.02705*, 2023.
- [29] Linyi Li, Jiawei Zhang, Tao Xie, and Bo Li. Double sampling randomized smoothing. In Kamalika Chaudhuri, Stefanie Jegelka, Le Song, Csaba Szepesvari, Gang Niu, and Sivan Sabato, editors, *Proceedings of the 39th International Conference on Machine Learning*, volume 162 of *Proceedings of Machine Learning Research*, pages 13163–13208. PMLR, 17–23 Jul 2022. URL <https://proceedings.mlr.press/v162/li22aa.html>.
- [30] Percy Liang, Rishi Bommasani, Tony Lee, Dimitris Tsipras, Dilara Soylu, Michihiro Yasunaga, Yian Zhang, Deepak Narayanan, Yuhuai Wu, Ananya Kumar, Benjamin Newman, Binhang Yuan, Bobby Yan, Ce Zhang, Christian Cosgrove, Christopher D Manning, Christopher Re, Diana Acosta-Navas, Drew A. Hudson, Eric Zelikman, Esin Durmus, Faisal Ladha, Frieda Rong, Hongyu Ren, Huaxiu Yao, Jue WANG, Keshav Santhanam, Laurel Orr, Lucia Zheng, Mert Yuksekgonul, Mirac Suzgun, Nathan Kim, Neel Guha, Niladri S. Chatterji, Omar Khattab, Peter Henderson, Qian Huang, Ryan Andrew Chi, Sang Michael Xie, Shibani Santurkar, Surya Ganguli, Tatsunori Hashimoto, Thomas Icard, Tianyi Zhang, Vishrav Chaudhary, William Wang, Xuechen Li, Yifan Mai, Yuhui Zhang, and Yuta Koreeda. Holistic evaluation of language models. *Transactions on Machine Learning Research*, 2023. ISSN 2835-8856. URL <https://openreview.net/forum?id=iO4LZibEqW>. Featured Certification, Expert Certification, Outstanding Certification.
- [31] Iman Mirzadeh, Keivan Alizadeh, Hooman Shahrokhi, Oncel Tuzel, Samy Bengio, and Mehrdad Farajtabar. Gsm-symbolic: Understanding the limitations of mathematical reasoning in large language models, 2024. URL <https://arxiv.org/abs/2410.05229>.
- [32] Mark Niklas Mueller, Franziska Eckert, Marc Fischer, and Martin Vechev. Certified training: Small boxes are all you need. In *The Eleventh International Conference on Learning Representations*, 2023. URL <https://openreview.net/forum?>

- 1079 [id=7oFuxtJtUMH](#).
- 1080 [33] David Noever. The enron corpus: Where the email bodies are buried?, 2020. URL <https://arxiv.org/abs/2001.10374>.
- 1081 [34] Alessandro De Palma, Harkirat S. Behl, Rudy R. Bunel, Philip H. S. Torr, and M. Pawan Kumar. Scaling the convex barrier with active sets. In *9th International Conference on Learning Representations, ICLR 2021, Virtual Event, Austria, May 3-7, 2021*, 2021.
- 1082 [35] Alessandro De Palma, Rudy R Bunel, Krishnamurthy Dj Dvijotham, M. Pawan Kumar, Robert Stanforth, and Alessio Lomuscio. Expressive losses for verified robustness via convex combinations. In *The Twelfth International Conference on Learning Representations*, 2024. URL <https://openreview.net/forum?id=mzyZ4wzKIM>.
- 1083 [36] Ethan Perez, Saffron Huang, Francis Song, Trevor Cai, Roman Ring, John Aslanides, Amelia Glaese, Nat McAleese, and Geoffrey Irving. Red teaming language models with language models, 2022. URL <https://arxiv.org/abs/2202.03286>.
- 1084 [37] Perplexity. Perplexity ai. AI Chatbot, 2023. URL <https://www.perplexity.ai/>.
- 1085 [38] Qwen, ., An Yang, Baosong Yang, Beichen Zhang, Binyuan Hui, Bo Zheng, Bowen Yu, Chengyuan Li, Dayiheng Liu, Fei Huang, Haoran Wei, Huan Lin, Jian Yang, Jianhong Tu, Jianwei Zhang, Jianxin Yang, Jiaxi Yang, Jingren Zhou, Junyang Lin, Kai Dang, Keming Lu, Keqin Bao, Kexin Yang, Le Yu, Mei Li, Mingfeng Xue, Pei Zhang, Qin Zhu, Rui Men, Runji Lin, Tianhao Li, Tianyi Tang, Tingyu Xia, Xingzhang Ren, Xuancheng Ren, Yang Fan, Yang Su, Yichang Zhang, Yu Wan, Yuqiong Liu, Zeyu Cui, Zhenru Zhang, and Zihan Qiu. Qwen2.5 technical report, 2025. URL <https://arxiv.org/abs/2412.15115>.
- 1086 [39] Gagandeep Singh and Deepika Chawla. Position: Formal methods are the principled foundation of safe AI. In *ICML Workshop on Technical AI Governance (TAIG)*, 2025. URL <https://openreview.net/forum?id=7V5CDSSjB7>.
- 1087 [40] Gagandeep Singh, Timon Gehr, Matthew Mirman, Markus Püschel, and Martin Vechev. Fast and effective robustness certification. *Advances in Neural Information Processing Systems*, 31, 2018.
- 1088 [41] Gagandeep Singh, Rupanshu Ganvir, Markus Püschel, and Martin Vechev. Beyond the single neuron convex barrier for neural network certification. In *Advances in Neural Information Processing Systems*, 2019.
- 1089 [42] Gagandeep Singh, Timon Gehr, Markus Püschel, and Martin Vechev. An abstract domain for certifying neural networks. *Proceedings of the ACM on Programming Languages*, 3(POPL), 2019.
- 1090 [43] Gagandeep Singh, Jacob Laurel, Sasa Misailovic, Debangshu Banerjee, Avaljot Singh, Changming Xu, Shubham Ugare, and Huan Zhang. Safety and trust in artificial intelligence with abstract interpretation. *Found. Trends Program. Lang.*, 8(3–4):250–408, June 2025. ISSN 2325-1107. doi: 10.1561/2500000062. URL <https://doi.org/10.1561/2500000062>.
- 1091 [44] Shubham Ugare, Debangshu Banerjee, Sasa Misailovic, and Gagandeep Singh. Incremental verification of neural networks. *Proc. ACM Program. Lang.*, 7(PLDI), June 2023. doi: 10.1145/3591299. URL <https://doi.org/10.1145/3591299>.
- 1092 [45] Shubham Ugare, Tarun Suresh, Debangshu Banerjee, Gagandeep Singh, and Sasa Misailovic. Incremental randomized smoothing certification. In *The Twelfth International Conference on Learning Representations*, 2024. URL <https://openreview.net/forum?id=SdeAPV1irk>.
- 1093 [46] Shubham Ugare, Rohan Gumaste, Tarun Suresh, Gagandeep Singh, and Sasa Misailovic. Itergen: Iterative semantic-aware structured llm generation with backtracking, 2025. URL <https://arxiv.org/abs/2410.07295>.
- 1094 [47] Caterina Urban and Antoine Miné. A review of formal methods applied to machine learning, 2021. URL <https://arxiv.org/abs/2104.02466>.
- 1095 [48] Boxin Wang, Weixin Chen, Hengzhi Pei, Chulin Xie, Mintong Kang, Chenhui Zhang, Chejian Xu, Zidi Xiong, Ritik Dutta, Rylan Schaeffer, Sang T. Truong, Simran Arora, Mantas Mazeika, Dan Hendrycks, Zinan Lin, Yu Cheng, Sanmi Koyejo, Dawn Song, and Bo Li. Decodingtrust: A comprehensive assessment of trustworthiness in gpt models, 2024. URL <https://arxiv.org/abs/2306.11698>.
- 1096 [49] Chengxiao Wang, Isha Chaudhary, Qian Hu, Weitong Ruan, Rahul Gupta, and Gagandeep Singh. Quantifying risks in multi-turn conversation with large language models, 2025. URL <https://arxiv.org/abs/2510.03969>.
- 1097 [50] Shiqi Wang, Kexin Pei, Justin Whitehouse, Junfeng Yang, and Suman Jana. Efficient formal safety analysis of neural networks. In *Advances in Neural Information Processing Systems*, 2018.
- 1098 [51] Shiqi Wang, Huan Zhang, Kaidi Xu, Xue Lin, Suman Jana, Cho-Jui Hsieh, and J Zico Kolter. Beta-CROWN: Efficient bound propagation with per-neuron split constraints for neural network robustness verification. In A. Beygelzimer, Y. Dauphin, P. Liang, and J. Wortman Vaughan, editors, *Advances in Neural Information Processing Systems*, 2021. URL <https://openreview.net/forum?id=ahYIIrBeCFw>.
- 1099 [52] Kaidi Xu, Zhouxing Shi, Huan Zhang, Yihan Wang, Kai-Wei Chang, Minlie Huang, Bhavya Kailkhura, Xue Lin, and Cho-Jui Hsieh. Automatic perturbation analysis for scalable certified robustness and beyond. In *Proceedings of the 34th International Conference on Neural Information Processing Systems, NIPS’20*, Red Hook, NY, USA, 2020. Curran Associates Inc. ISBN 9781713829546.
- 1100 [53] Kaidi Xu, Huan Zhang, Shiqi Wang, Yihan Wang, Suman Jana, Xue Lin, and Cho-Jui Hsieh. Fast and complete: Enabling complete neural network verification with rapid and massively parallel incomplete verifiers. In *International Conference on Learning Representations*, 2021. URL <https://openreview.net/forum?id=nVZtXBI6LNn>.

- 1128 [54] An Yang, Anfeng Li, Baosong Yang, Beichen Zhang, Binyuan Hui, Bo Zheng, Bowen Yu, Chang Gao, Chengan Huang,
1129 Chenxu Lv, Chujie Zheng, Dayiheng Liu, Fan Zhou, Fei Huang, Feng Hu, Hao Ge, Haoran Wei, Huan Lin, Jialong
1130 Tang, Jian Yang, Jianhong Tu, Jianwei Zhang, Jianxin Yang, Jiaxi Yang, Jing Zhou, Jingren Zhou, Junyang Lin, Kai
1131 Dang, Keqin Bao, Kexin Yang, Le Yu, Lianghao Deng, Mei Li, Mingfeng Xue, Mingze Li, Pei Zhang, Peng Wang,
1132 Qin Zhu, Rui Men, Ruize Gao, Shixuan Liu, Shuang Luo, Tianhao Li, Tianyi Tang, Wenbiao Yin, Xingzhang Ren,
1133 Xinyu Wang, Xinyu Zhang, Xuancheng Ren, Yang Fan, Yang Su, Yichang Zhang, Yingter Zhang, Yu Wan, Yuqiong
1134 Liu, Zekun Wang, Zeyu Cui, Zhenru Zhang, Zhipeng Zhou, and Zihan Qiu. Qwen3 technical report, 2025. URL
<https://arxiv.org/abs/2505.09388>.
- 1135 [55] Jehyeok Yeon, Isha Chaudhary, and Gagandeep Singh. Quantifying distributional robustness of agentic tool-selection,
1136 2025. URL <https://arxiv.org/abs/2510.03992>.
- 1137 [56] Huan Zhang, Tsui-Wei Weng, Pin-Yu Chen, Cho-Jui Hsieh, and Luca Daniel. Efficient neural network robustness
1138 certification with general activation functions. *Advances in neural information processing systems*, 31, 2018.
- 1139 [57] Huan Zhang, Shiqi Wang, Kaidi Xu, Linyi Li, Bo Li, Suman Jana, Cho-Jui Hsieh, and J Zico Kolter. General cutting
1140 planes for bound-propagation-based neural network verification. In Alice H. Oh, Alekh Agarwal, Danielle Belgrave,
1141 and Kyunghyun Cho, editors, *Advances in Neural Information Processing Systems*, 2022. URL <https://openreview.net/forum?id=5haAJAcofjc>.
- 1142 [58] Andy Zou, Zifan Maier, Daniel Liu, J Zachary Meng, Teodora Serrano, Matt Fredrikson, Pavol Mazeika, Jacob Stein-
1143 hardt, and Zico Kolter. Universal and transferable adversarial attacks on aligned language models. *arXiv preprint*
arXiv:2307.15043, 2023.
- 1144
- 1145
- 1146
- 1147
- 1148
- 1149
- 1150
- 1151
- 1152
- 1153
- 1154
- 1155
- 1156
- 1157
- 1158
- 1159
- 1160
- 1161
- 1162
- 1163
- 1164
- 1165
- 1166
- 1167
- 1168
- 1169
- 1170
- 1171
- 1172
- 1173
- 1174
- 1175
- 1176

1177 A Decoding Strategies

1178 Greedy decoding is a deterministic strategy that picks the highest probability next-token at each
 1179 step. Sampling-based methods sample the next token from a probability distribution modified
 1180 with parameters like *temperature*, top_p , top_k . Temperature smooths or sharpens the probability
 1181 distribution before sampling, top_p and top_k filter out low probability tokens from the probability
 1182 distribution. When sampling with temperature as $\tau \in (0, \infty)$

$$1184 P_M(x_i) = \sigma(z_i/\tau) = e^{z_i/\tau} / \sum_j e^{z_j/\tau}$$

1185
 1186 As $\tau \rightarrow 0$ sampling becomes more greedy and deterministic, whereas when $\tau \rightarrow \infty$ the probability
 1187 distribution approaches a uniform distribution. For top_k as $k \in \mathbb{N}$, let $V_k \subseteq V$ be the k tokens with
 1188 highest probability under P_M . Top $_k$ sampling restricts

$$1191 P_k(x_t | x_1 x_2 \dots x_{t-1}) = \begin{cases} P_M(x_t | x_1 x_2 \dots x_{t-1}) / \sum_{x' \in V_k} P_M(x' | x_1 x_2 \dots x_{t-1}) & x_t \in V_k \\ 0 & \text{otherwise} \end{cases}$$

1192
 1193 Similarly, for top_p (Nucleus sampling) [23] as $p \in (0, 1]$, let V_p be the minimal subset of V such
 1194 that $\sum_{x \in V_p} P_M(x | x_1 x_2 \dots x_{t-1}) \geq p$ where tokens in V_p are ordered by descending probability.

$$1197 P_p(x_t | x_1 x_2 \dots x_{t-1}) = \begin{cases} P_M(x_t | x_1 x_2 \dots x_{t-1}) / \sum_{x' \in V_p} P_M(x' | x_1 x_2 \dots x_{t-1}) & x_t \in V_p \\ 0 & \text{otherwise} \end{cases}$$

1200 B Rejection Sampling

1203 Algorithm 3: Rejection Sampling

1204 **Input** : Language Model M , Semantic Φ , Grammar G and Budget δ

1205 **Output**: P_{UB}, P_{LB}

```

1206 1  $P_{UB} \leftarrow 1.0$ ,  $P_{LB} \leftarrow 0.0$ ;
1207 2  $t \leftarrow 0$ ;
1208 3  $S \leftarrow \text{Set}()$  while  $t \leq \delta$  do
1209   4  $s, \mu(s) \leftarrow \text{Sample Sequence from Model } M$  ;
1210   5  $t \leftarrow t + |s|$ ;
1211   6 if  $s \notin S$  then
1212     7  $S \leftarrow S \cup \{s\}$ ;
1213     8 if  $s \models \Phi$  then
1214       9  $P_{LB} \leftarrow P_{LB} + \mu(s)$ ;
1215     10 else
1216       11  $P_{UB} \leftarrow P_{UB} - \mu(s)$ ;
1217     12 end
1218   13 end
1219 14 end
1220 15 return  $P_{UB}, P_{LB}$ 
1221
1222
1223
1224
1225
```

1226 C GSM-Symbolic Dataset

```

1227 You are an expert in solving grade school math tasks. You will be presented with a
1228 grade-school math word problem with symbolic variables and be asked to solve
1229 it.
1230 Only output the symbolic expression wrapped in <> that answers the question.
1231 The expression must use numbers as well as the variables defined in the
1232 question. You are only allowed to use the following operations: +, -, /, //,
1233 %, *, and **.
1234 You will always respond in the format described below: \n<>
1235 There are {t} trees in the {g}. {g} workers will plant trees in the {g} today.
1236 After they are done, there will be {tf} trees. How many trees did the {g}
1237 workers plant today?
<>tf - t>>
1238 If there are {c} cars in the parking lot and {nc} more cars arrive, how many cars
1239 are in the parking lot?
<>c + nc>>
1240 {p1} had {ch1} {o1} and {p2} had {ch2} {o1}. If they ate {a} {o1}, how many pieces
1241 do they have left in total?
<>ch1 + ch2 - a>>
1242 {p1} had {l1} {o1}. {p1} gave {g} {o1} to {p2}. How many {o1} does {p1} have left?
1243 <>l1 - g>>
1244 {p1} has {t} {o1}. For Christmas, {p1} got {tm} {o1} from {p2} and {td} {o1} from
1245 {p3}. How many {o1} does {p1} have now?
<>t + tm + td>>
1246 There were {c} {o1} in the {loc}. {nc} more {o1} were installed each day, from {d1}
1247 to {d2}. How many {o1} are now in the {loc}?
<>c + nc * (d2 - d1 + 1)>>
1248 {p1} had {gb1} {o1}. On {day1}, {p1} lost {l1} {o1}. On {day2}, {p1} lost {l2}
1249 more. How many {o1} does {p1} have at the end of {day2}?
<>gb1 - l1 - l2>>
1250 {p1} has ${m}. {p1} bought {q} {o1} for ${p} each. How much money does {p1} have
1251 left?
<>m - q * p>>
1252 {s2} has a bag of {s3} with {d} inside. He tripped over {s4} while carrying it and
1253 dropped {b} of them. He scrambled to search for them but only came up with {c}.
1254 When he went back home, he inspected the {s3} further. {a} of them he
1255 picked up weren't {s3}, but actually {s1} so he got rid of it. How many {s3}
1256 did {s2} end up with?

```

1260 Listing 1. Example prompt for the GSM-Symbolic task [31].
1261

```

1262 start: SPACE? "<<" SPACE? expr SPACE? ">>" SPACE?
1263 expr: term (SPACE? ("+" | "-") SPACE? term)*
1264 term: factor (SPACE? ("*" | "//" | "/" | "%") SPACE? factor)*
1265 factor: "-" SPACE? factor
1266     | TYPE "(" SPACE? expr SPACE? ")"
1267     | primary SPACE?
1268 primary: NUMBER
1269     | VARIABLE
1270     | "(" SPACE? expr SPACE? ")"
1271 TYPE: "int"
1272 SPACE: " "
1273 DIGIT: /[0-9]/
1274 INT: DIGIT+
1275 SIGNED_INT: (( "+" | "-"))? INT

```

```

1275 DECIMAL: INT "." INT?
1276   | "." INT
1277 EXP: ("e" | "E") SIGNED_INT
1278
1279 FLOAT: INT EXP
1280   | DECIMAL EXP?
1281
1282 NUMBER: FLOAT
1283   | INT
1284
1285 LCASE LETTER: /[a-z]/
1286 UCASE LETTER: /[A-Z]/
1287
1288 LETTER: LCASE LETTER
1289   | UCASE LETTER
1290
1291 CNAME: ("_" | LETTER) ("_" | LETTER | DIGIT))*
```

1290 Listing 2. GSM-Symbolic Grammar [31]

1291

D Email Leakage Dataset

1292

```

1293 Please continue the below context by ONLY outputting the email address that
1294 continues the context:
1295 the email address of Sean P. Tracey is seanpat@flash.net; the email address of
1296 Hershel Swartz is hershel@flash.net; the email address of James Shirley is
1297 jrshirly@flash.net; the email address of Patrick Locke is cannon2@flash.net;
the email address of Mom is bmenconi@flash.net; the email address of Karen
1298 Arnold is
```

1299

1300 Listing 3. Example prompt for the Enron Email Leakage Task task [33].

1301

1302

1303

1304

1305

1306

1307

1308

1309

1310

1311

1312

1313

1314

1315

1316

1317

1318

1319

1320

1321

1322

1323